

# What *Not* to Do When Writing an Interpreter for Specialisation

Neil D. Jones\*

DIKU, University of Copenhagen  
e-mail: `neil@diku.dk`

**Abstract.** A partial evaluator, given a program and a known “static” part of its input data, outputs a *specialised* or *residual* program in which computations depending only on the static data have been performed in advance.

*Ideally* the partial evaluator would be a “black box” able to extract nontrivial static computations whenever possible; which never fails to terminate; and which always produces residual programs of reasonable size and maximal efficiency, so all possible static computations have been done. *Practically* speaking, partial evaluators often fall short of this goal; they sometimes loop, sometimes pessimise, and can explode code size.

A partial evaluator is analogous to a spirited horse: while impressive results can be obtained when used well, the user must know what he/she is doing. Our thesis is that *this knowledge can be communicated* to new users of these tools.

This paper presents a series of examples, concentrating on a quite broad and on the whole quite successful application area: using specialisation to remove interpretational overhead. It presents both positive and negative examples, to illustrate the effects of program style on the efficiency and size of the target programs obtained by specialising an interpreter with respect to a known source program. It concludes with a checklist summarising what was learned from the examples, discussions, and problem analyses.

## 1 Introduction

### 1.1 The phenomena

In my opinion (and, I hope, in yours too after looking at the following examples), there exist a number of *real phenomena* in specialisation, i.e. frequently occurring specialisation problems that are nontrivial to solve; there exist styles of programming that can be expected to give *any* partial evaluator difficulties; and one can *learn and communicate* programming style that leads to good specialisation, i.e. high speedups without explosion of code size.

---

\* This research was partially supported by the Danish *DART* project, and the Esprit BRA *Atlantique*.

## 1.2 Making implicit knowledge explicit

Some of the following examples will make the experienced reader grit his or her teeth, saying “I would never dream of programming my interpreter in such an absurd way!” This is indeed true — but inexperienced users often commit just these errors of thought or of programming style (I speak from experience with students.) A major point of this article is to make manifest some of the *implicit knowledge* we have acquired in our field, so newcomers can avoid encountering and having to figure out how to solve the same problems in “the school of hard knocks.”

## 1.3 Dependence on particular specialisation algorithms

The advice alluded to in this paper’s title would seem, at least to some extent, relative to which *class of specialisation algorithms* is used. I have carefully written to avoid having to take account of remarks of the form “yes, but I can solve that problem using feature XXX of my partial evaluator.”

Accounting for such possible remarks, followed by responses (“even so, ...”), etc. could easily tend to lead to an infinite regress, or to a protracted discussion whose line would be hard to follow. Even worse, it would make it hard at the end to say what really had been accomplished by the discussion.

## 2 Concepts and notations

We use familiar notations, for example  $D$  is a data set that contains all first-order values, including *all program texts*. Inputs may be paired, e.g. “ $\mathbf{d.e}$ ” is in  $D$  if both  $\mathbf{d}$  and  $\mathbf{e}$  are in  $D$ . We may also write  $(\mathbf{d.e})$  or  $(\mathbf{d}, \mathbf{e})$ . For concreteness the reader may think of  $D$  as Lisp/Scheme lists, and programs as written in Scheme. However the points below are not language-dependent, and some example interpreters will be written in imperative languages.

The *semantic function* of program  $\mathbf{p}$  is a partial function  $\llbracket \mathbf{p} \rrbracket : D \rightarrow D \cup \{\perp\}$ , so  $\llbracket \mathbf{p} \rrbracket(\mathbf{d}) = \mathbf{d}' \in D$  is the result of running  $\mathbf{p}$  on input  $\mathbf{d} \in D$  (and is  $\perp$  if  $\mathbf{p}$  fails to terminate on  $\mathbf{d}$ .)

We will assume the underlying implementation language is defined by an (unspecified) operational semantics. Computations are given by *sequences* (for imperative languages) or, more generally, by *computation trees*.

Further, program run time is assumed roughly proportional to the size of the computation tree deducing  $\llbracket \mathbf{p} \rrbracket(\mathbf{d}) = \mathbf{d}'$ . (This will become relevant in the discussion of linear speedup.) Notation:  $\text{time}_{\mathbf{p}}(\mathbf{d}) \in \mathbb{N}$  is the time taken to compute  $\llbracket \mathbf{p} \rrbracket(\mathbf{d})$  (and is  $\infty$  if  $\mathbf{p}$  fails to terminate on  $\mathbf{d}$ .)

Given program  $\mathbf{p}$  and “static” input  $\mathbf{s} \in D$ , program  $\mathbf{p}_{\mathbf{s}}$  is the *result of specialising  $\mathbf{p}$  to  $\mathbf{s}$*  if  $\llbracket \mathbf{p} \rrbracket(\mathbf{s.d}) = \llbracket \mathbf{p}_{\mathbf{s}} \rrbracket(\mathbf{d})$  for every “dynamic” input  $\mathbf{d} \in D$ . Well-known terms are to call  $\mathbf{p}_{\mathbf{s}}$  a *residual program* for  $\mathbf{p}$  with respect to  $\mathbf{s}$ , or a version

of  $p$  specialised to  $s^2$ .

A *partial evaluator* (or *specialiser*) is a program  $\text{mix}$  such that for every program  $p$  and “static” input  $s \in D$ , program  $\llbracket \text{mix} \rrbracket(p.s)$  is a version of  $p$  specialised to  $s$ .

Clearly a trivial  $\text{mix}$  is easy to build (e.g. by Kleene’s *s-m-n* Theorem of the 1930s.) The practical goal is to make  $p_s = \llbracket \text{mix} \rrbracket(p.s)$  *as efficient as possible*, by performing many of  $p$ ’s computations — ideally, all those that depend on  $s$  alone. The *speedup* realised by  $\text{mix}$  on program  $p$  and static input  $s$  is the following ratio:

$$\text{speedup}_s(d) = \frac{\text{time}_p(s, d)}{\text{time}_{p_s}(d)}$$

(Note that for each static input  $s$  it is a function of  $d$ .)

Some programs are well-suited to specialisation, yielding substantial speedups, and others are not. Given a specialiser  $\text{mix}$ , a *binding-time improvement* is the transformation of program  $p$  into an equivalent program  $q$  such that  $\llbracket p \rrbracket = \llbracket q \rrbracket$ , but specialisation of  $q$  to various  $s$  gives larger speedups than specialisation of  $p$  to the same  $s$ .

*Other languages.* If unspecified, as above, we assume a standard implementation language  $L$  is intended. Otherwise, if another language  $S$  is used, we make this explicit as follows. The *semantic function* of a program  $p$  is a partial function  $\llbracket p \rrbracket^s : D \rightarrow D \cup \{\perp\}$ , so  $\llbracket p \rrbracket^s(d)$  is the result of running  $p$  on input  $d \in D$  ( $\perp$  if  $p$  fails to terminate on  $d$ .) Further,  $\text{time}_p^s(d)$  is the time taken to compute  $\llbracket p \rrbracket^s(d)$  (and is  $\infty$  if  $p$  fails to terminate on  $d$ .)

**Interpreters.** Program  $\text{int}$  is an *interpreter for language S* if for all  $S$ -programs  $s$  and data  $d \in D$

$$\llbracket \text{int} \rrbracket(s.d) = \llbracket s \rrbracket^s(d)$$

### 3 Assumptions in this paper

For definiteness and to avoid having to “hit a moving target,” I now list the assumptions of this paper. Readers are welcome to reply by describing how other techniques can solve some or all of these problems; but should be honest enough to admit it if their techniques introduce yet other problems.

---

<sup>2</sup> This definition concerns only correctness of specialisation, and takes no account of how specialisation is done. It is thus a purely *extensional* definition, ignoring “intensional” features of  $p$  such as efficiency or size; these are treated for concrete examples in the following text. A more intensional view is often taken in logic programming [7], and sometimes by other authors in functional and imperative programming.

### 3.1 Restrictions in this discussion

1. The programs to be specialised are assumed expressed in an *untyped first-order functional* language, using an informal syntax. Consequently all our interpreter texts, etc. are in this form; however the *languages they interpret* may be functional, imperative, higher-order, logic, etc.
2. Offline specialisation is assumed: before actual specialisation begins, all of program **p**'s operations and calls or other control transfers are *annotated* as either **static**: do at specialisation time; or as **dynamic**: generate code to be executed at run time.
3. Every function parameter (in the program to be specialised) is either *totally static* or *totally dynamic*. Thus there is no *partially static data* in these examples.
4. We assume that a BTA (*binding-time analysis*) does the annotations, e.g. separately as in Schism, or integrated into the text as in Similix:

```
append(X_S,Y_D) = if_S X=_S 'nil then Y
                  else cons_S(hd_S(X), append_S(tl_S X, Y))
```

5. No *distributivity rules* are used, e.g. to transform from 1 to 2 below:

```
1: static1 + (if dynamic_test then static2 else static3)
2: if dynamic_test then (static1+static2) else (static1+static3)
```

This has been called “CPS specialisation” — a rather unfortunate name in my opinion since the concept makes perfect sense without continuations, or even without higher-order functions. For instance, distributivity or “CPS specialisation” can even be done in logic programming:

```
p(D, Result) :- test(D, S), Result is S1 + S.
test([], S2).
test([X|Xs], S3).
```

is equivalent to

```
p([], Result)      :- Result is S1 + S2.
p([X|Xs], Result) :- Result is S1 + S3.
```

6. *The Trick* or  $\eta$ -conversion will not be used unless explicitly stated. This is a programming device (described in [6]) to make static otherwise dynamic values, provided they are known to be of bounded static variation (this term is defined below.)
7. No global optimisations are done, e.g.

**dead code elimination** (a logic programming analog is removing code that is certain to fail); or

**common subexpression elimination**, especially ones containing function calls (a logic programming analog is removing duplicated goals.)

### 3.2 Remarks on the restrictions

The lack of partially static data can be overcome in more than one way; for instance it is of less significance for online specialisers such as Fuse [13], super-compilation [12], or most logic programming specialisers [10, 4]; but then new problems arise concerning control, termination, etc.

An offline specialiser such as Similix [1] or Schism [2] usually does analyses, often involving some form of tree grammar [8], to statically describe all possible runtime data shapes and to exploit this. Alternatively, a *bifurcation* transformation [9, 3] splits a program with partially static data into one part with completely static data, and another part with both static and dynamic data.

*Online and offline specialisation.* I maintain that many of specialisation *problems* are essentially independent of which type of specialisation *algorithm* one uses, or of whether one is specialising functional, logic, or imperative programs. This holds even though a particular problem might manifest itself in different ways, e.g. in an offline functional specialiser as a *nontermination problem*, or in logic programming as causing an *unnecessarily conservative* specialisation, leading to little or no speedup. Further, I claim that:

- Several of the “binding-time improvements” exemplified below to avoid non-termination or code size explosion would improve specialisation of logic programs as well as functional programs.
- Binding-time improvements done to improve offline specialiser results often improve online specialisation results too. This can either be by causing more to be statically computable, or by reducing the size of the residual program.

## 4 Bounded static variation: the key to termination of the specialiser

Nontermination at specialisation time can occur for one of two reasons: an

- Attempt to build an *infinitely large* command, expression or literal; or an
- Attempt to build a residual program containing *infinitely many* program points (labels, defined function, or procedures.)

The basic cause is the same: the specialiser’s *unfolding strategy* is the main cause of the one or the other misbehaviour. In either case nontermination makes a specialiser less suitable for use by nonexperts, and *quite unsuitable for automatic use*. Further, failure to terminate gives the user very little to go on to discover the problem’s cause, in order to repair it.

Two attitudes to nontermination have been seen: in functional programming, infinite static loops are often seen as the user’s fault, so there is no need for the specialiser to account for them. Logic programming most often requires that the specialiser *always terminate*. One reason is the Prolog “negation as finite failure” semantics, which means that changing termination properties can change

semantics, and even answer sets. Each view can be (and has been) both defended and attacked.

Current work on functional program specialisation concerns a BTA conservative enough to guarantee that specialisation always terminates, and still gives good results when specialising interpreters [5].

#### 4.1 Pachinko, execution, and specialisation

Normal execution follows only one sequential control thread. Specialisation, however, must account for *all possible control threads* the subject program could enter into, for all possible dynamic inputs.

An analogy: the popular Japanese “Pachinko” entertainment involves steel balls that fall through a lattice of pins. Sequential execution corresponds to dropping only one ball, which thus traverses only one trajectory. On the other hand, specialisation to static input  $\mathbf{s}$  corresponds to dropping an infinite set of balls at once — all those which share  $\mathbf{s}$  as given static input, but have *all possible* dynamic input values. Reasoning about specialisation, e.g. residual program finiteness or efficiency, requires reasoning about such a (most often infinite) set of trajectories.

*Reachable states.* Let function<sup>3</sup>  $\mathbf{f}$  defined in program  $\mathbf{p}$  have number of parameters denoted by “arity( $\mathbf{f}$ )” and let parameter values range over the set  $V$ . The  $i$ th parameter of  $\mathbf{f}$  will be written as  $\mathbf{f}_i$ .

A *state* is a pair  $(\mathbf{f}, \bar{\alpha})$  where  $\bar{\alpha} \in V^{\text{arity}(\mathbf{f})}$ . Notation:  $\bar{\alpha}_i$  is the  $i$ th component of parameter tuple  $\bar{\alpha} \in V^{\text{arity}(\mathbf{f})}$ . Thus  $\bar{\alpha}_i$  is the value of  $\mathbf{f}_i$  in state  $(\mathbf{f}, \bar{\alpha})$ .

State  $(\mathbf{f}, \bar{\alpha})$  can *reach* state  $(\mathbf{g}, \bar{\beta})$ , written  $(\mathbf{f}, \bar{\alpha}) \rightarrow (\mathbf{g}, \bar{\beta})$ , if computation of  $\mathbf{f}$  on parameter values  $\bar{\alpha}$  directly requires the value of  $\mathbf{g}$  on parameter values  $\bar{\beta}$ . (Formal definition of this is straightforward.)

#### 4.2 Definition of bounded static variation.

Let ‘;’ be the parameter *tuple concatenation operator*, suppose  $\mathbf{f1}$  is the initial (entry) function<sup>3</sup> of program  $\mathbf{p}$ , and let  $s, d \in \mathbb{N}$ , where  $s + d = \text{arity}(\mathbf{f1})$ . We assume (for simplicity) that the first  $s$  parameters of  $\mathbf{f1}$  are static, and the remaining  $d$  parameters are dynamic. The set of *statically reachable* states  $\mathcal{SR}(\bar{\sigma})$  consists of all states  $(\mathbf{f}, \bar{\alpha})$  reachable in 0, 1, or more steps from some initial call with  $\bar{\sigma} \in V^s$  as static program input, and an arbitrary dynamic input  $\bar{\delta}$ :

$$\mathcal{SR}(\bar{\sigma}) = \{(\mathbf{f}, \bar{\alpha}) \mid \exists \bar{\delta} \in V^d : (\mathbf{f1}, \bar{\sigma}; \bar{\delta}) \rightarrow^* (\mathbf{f}, \bar{\alpha})\}$$

The  $i$ th parameter  $\mathbf{f}_i$  of  $\mathbf{f}$  is of *bounded static variation*, written  $\text{BSV}(\mathbf{f}_i)$ , iff for all  $\bar{\sigma} \in V^s$ , the following set is finite:

$$\{\bar{\alpha}_i \mid (\mathbf{f}, \bar{\alpha}) \in \mathcal{SR}(\bar{\sigma})\}$$

---

<sup>3</sup> These definitions are easy to extend to imperative languages, and are not difficult to extend to logic programming.

This captures the intuitive notion that  $f$ 's  $i$ th parameter only varies finitely during all computations in which the static program input is fixed to any  $\bar{\sigma}$ , and dynamic program input varies freely.

Such a parameter may, in principle, be classified as “static.” Most such parameters are directly computable from the static program inputs. Exceptions do occur, for example the following expression is of BSV even if  $x$  is dynamic:

```
if x mod 2 = 0 then 'EVEN else 'ODD
```

## 5 Compositionality and semicompositionality of interpreters

Suppose each of an interpreter's function parameters has been classified as either nonsyntactic, or as *syntactic*, in which case its values range only over phrases from the source language — perhaps or perhaps not extracted from the interpreted source program, henceforth called **src**.

The interpreter is defined to be **compositional** if for any definition of a function with syntactic parameters, the syntactic parameters of any subcall involving recursion are always *proper substructures of the calling function's* syntactic parameters. Clearly in any computation by a compositional interpreter, the syntactic parameters of any functions called are necessarily substructures of the source program **src** being interpreted. Thus they are quite obviously of BSV.

The compositionality requirement ensures that the interpreter can be specialised with respect to its static input **src** by simple unfolding — a process *guaranteed to terminate* (though perhaps not fully satisfactory, unless additional specialisation-time computations are done.)

A weaker requirement, also guaranteeing BSV, is for the interpreter to be **semicompositional**, meaning that called parameters must be *substructures of the original source program src*, but need not decrease locally in every call.

Semicompositionality can even allow syntactic parameters to grow; the only requirement is that the set of their values is limited to range over the set of all substructures of **src** (perhaps including the whole of **src** itself.) Consequently, **while**'s meaning can be defined in terms of itself, and procedure calls may be elaborated by applying an execution function to the body of the called procedure (both of which violate strict compositionality.)

Semicompositionality does not guarantee termination at run time; but it does guarantee finiteness of the set of specialised program points — so *termination of specialisation* can be ensured by a natural unfolding strategy (unfold, unless this function has been seen before with just these static parameter values.)

## 6 Where do interpreters come from?

An interpreter is just a program. It can be used to define a new language (so a source program meaning is whatever the interpreter yields when given the source program and *its* input); or to implement a language already defined by other means, e.g. by an operational or denotational semantics.

### Rules for Expressions

Constant :	$\langle S, M, c \cdot C \rangle$	$\Rightarrow \langle c \cdot S, M, C \rangle$
Variable :	$\langle S, M, xi \cdot C \rangle$	$\Rightarrow \langle M(xi) \cdot S, M, C \rangle$
Composite :	$\langle S, M, (e_1 \text{ op } e_2) \cdot C \rangle$	$\Rightarrow \langle S, M, e_1 \cdot e_2 \cdot \text{op} \cdot C \rangle$
Operator :	$\langle \underline{n_2} \cdot \underline{n_1} \cdot S, M, \text{op} \cdot C \rangle$	$\Rightarrow \langle \underline{n_1 \text{ op } n_2} \cdot S, M, C \rangle$

### Rules for Programs

Null :	$\langle S, M, \text{null} \cdot C \rangle$	$\Rightarrow \langle S, M, C \rangle$
Assign :	$\langle S, M, (xi := e) \cdot C \rangle$	$\Rightarrow \langle i \cdot S, M, e \cdot \text{assign} \cdot C \rangle$
Sequence :	$\langle S, M, (p_1 ; p_2) \cdot C \rangle$	$\Rightarrow \langle S, M, p_1 \cdot p_2 \cdot C \rangle$
Test :	$\langle S, M, (\text{if } b \text{ then } p_1 \text{ else } p_2) \cdot C \rangle$	$\Rightarrow \langle p_1 \cdot p_2 \cdot S, M, b \cdot \text{if} \cdot C \rangle$
Loop :	$\langle S, M, (\text{while } b \text{ do } p_1) \cdot C \rangle$	$\Rightarrow \langle b \cdot p_1 \cdot S, M, b \cdot \text{while} \cdot C \rangle$

### Rules for assign, if and while

$\langle \underline{n} \cdot i \cdot S, M, \text{assign} \cdot C \rangle$	$\Rightarrow \langle S, M[xi \mapsto \underline{n}], C \rangle$
$\langle \text{true} \cdot p_1 \cdot p_2 \cdot S, M, \text{if} \cdot C \rangle$	$\Rightarrow \langle S, M, p_1 \cdot C \rangle$
$\langle \text{false} \cdot p_1 \cdot p_2 \cdot S, M, \text{if} \cdot C \rangle$	$\Rightarrow \langle S, M, p_2 \cdot C \rangle$
$\langle \text{true} \cdot b \cdot p_1 \cdot S, M, \text{while} \cdot C \rangle$	$\Rightarrow \langle S, M, p_1 \cdot (\text{while } b \text{ do } p_1) \cdot C \rangle$
$\langle \text{false} \cdot b \cdot p_1 \cdot S, M, \text{while} \cdot C \rangle$	$\Rightarrow \langle S, M, C \rangle$

Figure 6.1: Instructional interpreter `int-instruct`.

## 6.1 Interpreters derived from execution models

Many interpreters simply embody some model of runtime execution, e.g. Pascal's "stack of activation records" or Algol 60's "thunk" mechanism. Following is a simple example.

### An interpreter used for instruction at DIKU.

Computation is by a linear sequence

$$(S_1, M_1, C_1) \rightarrow (S_2, M_2, C_2) \rightarrow \dots$$

of states of form  $(S, M, C)$ .  $S$  is a *computation stack*,  $M : \{x_0, x_1, \dots\} \rightarrow \text{Value}$  is a *memory* containing the values of variables  $xi$ , and  $C$  is a *control stack*, containing bits of commands and expressions that remain to be evaluated/executed. (Remark:  $C$  is in essence a materialisation in the form of data of the program's *continuation* from the current point of execution.)

For readability we describe this interpreter (call it `int-instruct`) by a set of transition rules in Figure 6.1. Note that the  $C$  component is neither compositional nor semicompositional. The underlines, e.g.  $\underline{n}$ , indicate numeric runtime values, and  $M[xi \mapsto \underline{n}]$  is a memory  $M'$  identical to  $M$ , except that  $M'(xi) = \underline{n}$ .



## 6.2 Interpreters derived from denotational semantics

A *denotational semantics* [11] consists of

- a collection of *domain equations* specifying the partially ordered value sets to which the meanings (denotations) of program pieces and auxiliaries such as environments will belong, and
- a collection of *semantic functions* mapping program pieces to their meanings in a compositional way (as defined in Section 5.)

In practice a denotational semantics resembles a program in a modern functional language such as ML, Haskell, or Scheme. Indeed the designers of such languages have to a certain extent built denotational concepts into them. One instance is syntax and implementation techniques to make programs with higher-order functions easy to write and not too expensive to use. Another notable instance is *continuations*, which have come from the metalevel of language descriptions (an invention devised to formalise **goto**) down to the subject language level.

A difference from operational definitions is that fixpoints are conceptually evaluated “all at once” in denotational semantics, and only unfolded “upon demand” in functional languages. Still, the net effect is that many denotational language definitions can simply be transliterated into functional programs.

One problem with a denotational semantics is that it necessarily uses a “universal domain,” capable of expressing *all possible run-time values* of all programs that can be given meaning. This, combined with the requirement of compositionality, often leads to an (over-)abundance of higher-order functions. These two together imply that domain definitions very often must be both *recursive and higher-order*, creating both mathematical and implementational complications (e.g. reflexive domains and contravariance in a function space functor.)

Examples are familiar, and omitted for brevity.

## 6.3 Interpreters derived from operational semantics

*Operational semantics* seems more manageable, since emphasis is on judgements formed from first-order objects (although finite functions are allowed.) They come in two flavours: *big-step* in which a judgement may, for example, be of form

$$\dots \vdash Exp \Rightarrow \text{Final answer}$$

and *small-step* in which a judgement is of form

$$\dots \vdash Exp \Rightarrow Exp'$$

which describes expression evaluation by reducing one expression repeatedly to another, until a final answer is obtained. In both cases a language definition is a set of *inference rules*. A judgement is proven to hold by constructing a proof tree with it as root, and where each local subtree is an instance of an inference

```

; Type is: Program x Program_value -> Output_value

Run(pgm, input) = Eval(e1, ns1, cons(input, 'nil), pgm) where
    e1 = lookbody(first_fcn(pgm), pgm)
    ns1 = lookfcns(first_fcn(pgm), pgm)

; Type is: Expression x Namelist x Valuelist x Program ->
Value

Eval(e, ns, vs, pgm) =    case e of
    constant            : constant
    'X                  : lookparams(X, ns, vs)

    '(e1 binop e2)      : apply(binop, v1, v2) where
                        v1 = Eval(e1, ns, vs, pgm)
                        v2 = Eval(e2, ns, vs, pgm)

    '(if e0 e1 e2)      : if Eval(e0, ns, vs, pgm)
                        then Eval(e1, ns, vs, pgm)
                        else Eval(e2, ns, vs, pgm)

    '(call f es)        : Eval(e1, ns1, vs1, pgm) where
                        e1 = lookbody(f, pgm)
                        ns1 = lookfcns(f, pgm)
                        vs1 = Evlist(es, ns, vs, pgm)

Evlist(es, ns, vs, pgm) =    case es of
    'nil                : nil
    'cons(e1 es1)       : cons(Eval(e1, ns, vs, pgm),
                                Evlist(es1, ns, vs, pgm))

```

Figure 6.2: *Big-step semantics int-big-step in program form.*

rule. Examples are omitted for brevity, but the approaches should be clear from the following program implementation examples.

Transliterating these two types of operational semantics into program form gives interpreters of rather different characteristics. Big-step semantics are conceptually nearer denotational semantics than small-step semantics, since source syntax and program meanings are clearly separated. Small-step semantics work by symbolic source-to-source code reduction, and so are nearer equational or rewrite theories, for example traditional treatments of the  $\lambda$ -calculus.

**An interpreter derived from a big-step operational semantics.** The interpreter `int-big-step` of Figure 6.2 implements a “big-step” semantics. The program has only one input, but its internal functions may have any number of parameters. This may be seen in the interpretation of `(call f es)`.

**Interpreters derived from small-step operational semantics.**

At first sight, a small-step operational semantics seems quite unsuitable for specialisation. A familiar example is the usual definition of reducibility in the lambda calculus. This consists of a few rules for reduction, e.g.  $\alpha$ ,  $\beta$ ,  $\eta$ , and *context rules* stating that these can be applied when occurring inside other syntactic constructs. Problems with direct implementation include:

- *Nondeterminism*: many redexes may be reducible within the same  $\lambda$ -expression;
- *Nondirectionality of computation*: rewriting may occur either from left to right or vice versa (conversion versus reduction); and
- *Syntax rewriting*: the subject program’s syntax is continually changed, in ways impossible statically to predict while specialising an interpreter.

These are also problems for implementations, often resolved as follows:

- Nondeterminism: reduction is sometimes restricted to occur only from the outermost syntactic operator, using call-by-value, call-by-name, or some other fixed strategy;
- Nondirectionality: rewriting occurs only from left to right, until a normal form is obtained (one that cannot be rewritten further).

More generally, nondeterminism is often resolved by defining explicit *evaluation contexts*  $C[]$ . Each is an “expression with a hole in it” indicated by the  $[]$ . Determinism may be achieved by defining evaluation contexts so the *unique decomposition property* holds: any expression  $E$  has at most one decomposition of form  $E = C[E']$  where  $C[]$  is an evaluation context. Given this, computation proceeds through a series of rewritings of such configurations.

The problem of *syntax rewriting* causes particular problems for specialisation. The reason is that rewriting often leads to an infinity of different specialisation-time values. For example, this can occur in the  $\lambda$ -calculus if  $\beta$ -reduction is implemented by substitution in an expression with recursion, either implicitly by the  $Y$  combinator or explicitly by a **fix** construct.

The problem may be alleviated by using *closures* instead of rewriting, essentially a “lazy” form of  $\beta$ -reduction. This easily allows building a semicompositional big-step interpreter, but problems still remain for small-step interpreters, for example dealing with so-called “delta rules.”

Some small-step semantics have a property that can be thought of as *dual to semicompositionality*: every decomposition  $E = C[E']$  that occurs in a given computation is such that  $C[]$  consists of *the original program with a hole in it*, and  $E'$  is a normal form value, e.g. a number or a closure.

Such a dual semicompositional semantics can usually be specialised well. The condition is violated, however (for an example,) if the semantics implements function or procedure calls by substituting the body of the called function or procedure in place of the call itself (a similar example is seen in Section 7.2 below.)

### Rules for Expressions

Constant :  $\langle S, M, c \cdot C \rangle \Rightarrow \langle c \cdot S, M, C \rangle$   
 Variable :  $\langle S, M, xi \cdot C \rangle \Rightarrow \langle M(xi) \cdot S, M, C \rangle$   
 Composite :  $\langle S, M, (e_1 \text{ op } e_2) \cdot C \rangle \Rightarrow \langle S, M, e_1 \cdot e_2 \cdot op \cdot C \rangle$   
 Operator :  $\langle \underline{n_2} \cdot \underline{n_1} \cdot S, M, op \cdot C \rangle \Rightarrow \langle \underline{n_1 \text{ op } n_2} \cdot S, M, C \rangle$

### Rules for Programs

Null :  $\langle S, M, \text{null} \cdot C \rangle \Rightarrow \langle S, M, C \rangle$   
 Assign :  $\langle S, M, (xi := e) \cdot C \rangle \Rightarrow \langle S, M, e \cdot \text{assign} \cdot i \cdot C \rangle$   
 Seq. :  $\langle S, M, (p_1 ; p_2) \cdot C \rangle \Rightarrow \langle S, M, p_1 \cdot p_2 \cdot C \rangle$   
 Test :  $\langle S, M, (\text{if } b \text{ then } p_1 \text{ else } p_2) \cdot C \rangle \Rightarrow \langle S, M, b \cdot \text{if} \cdot p_1 \cdot p_2 \cdot C \rangle$   
 Loop :  $\langle S, M, p \cdot C \rangle \text{ if } p = (\text{while } b \text{ do } p_1) \Rightarrow \langle S, M, b \cdot \text{while} \cdot p_1 \cdot p \cdot C \rangle$

### Rules for assign, if and while

$\langle \underline{n} \cdot S, M, \text{assign} \cdot i \cdot C \rangle \Rightarrow \langle S, M[xi \mapsto \underline{n}], C \rangle$   
 $\langle \text{true} \cdot S, M, \text{if} \cdot p_1 \cdot p_2 \cdot C \rangle \Rightarrow \langle S, M, p_1 \cdot C \rangle$   
 $\langle \text{false} \cdot S, M, \text{if} \cdot p_1 \cdot p_2 \cdot C \rangle \Rightarrow \langle S, M, p_2 \cdot C \rangle$   
 $\langle \text{true} \cdot S, M, \text{while} \cdot p_1 \cdot p \cdot C \rangle \Rightarrow \langle S, M, p_1 \cdot p \cdot C \rangle$   
 $\langle \text{false} \cdot S, M, \text{while} \cdot p_1 \cdot p \cdot C \rangle \Rightarrow \langle S, M, C \rangle$

Figure 7.1: Modified instructional interpreter.

## 7 Bounded or unbounded static variation?

### 7.1 Analysis of the instructional interpreter.

The interpreter `int-instruct` of Figure 6.1 works quite well for computation and proof, and has been used for many exercises and proofs of program properties in the course “Introduction to Semantics.” On the other hand, it does not specialise at all well! The source of the problem: suppose one is given the source program  $p$ , i.e. the initial control stack contents are  $C = p \cdot ()$ , but that the memory  $M$  is unknown, i.e. dynamic.

Clearly the value stack  $S$  must also be dynamic, since it receives values from  $M$  as in the second transition rule. But this implies that *anything* put into  $S$  and then retrieved from it again must also be dynamic. In particular this includes: the index  $i$  of a variable to be assigned; the **then** and **else** branches of an **if** statement; and the test of a **while** statement. Consequently *all of these* become dynamic, and so essentially no specialisation at all occurs.

### A better version for specialisation.

The problem is easy to fix: just place the above-mentioned syntactic entities on the *control stack*  $C$  instead of  $S$ : the index  $i$  of a variable to be assigned; the **then** and **else** branches of an **if** statement; and the test of a **while** statement. This leads to the interpreter of Figure 7.1, again as a set of transition rules.

*Bounded static variation of the control stack.* In this version stack  $C$  is built up only from pieces of the original program. Further, given any one source program  $C$  takes on only finitely many different possible values. Even though  $C$  can grow when a **while** command is processed, it is easy to see that it cannot grow unboundedly for any one, fixed, source program. Thus  $C$  may safely be annotated as “static.”

This version specialises much better. The effect is to yield a target program consisting of labeled transitions from one  $(S, M)$  pair to another — in essence with one transition (more or less) for each point in the original program, and no source code at all. Many of these transitions involve no change either to  $S$  or  $M$ , and so are removable by a trivial “transition compression.”

## 7.2 A small but problematic extension: procedures.

Procedures are easily added by adding a fourth component  $(S, M, C, Pdefs)$  where  $Pdefs$  is a list of mutually recursive procedure definitions. A procedure call is handled by looking up the procedure’s name in  $Pdefs$ , and replacing the call by the body of the called procedure.

### Rule for Procedure declarations

$$\langle S, M, (\text{procedure } P : p) \cdot C, Pdefs \rangle \Rightarrow \langle S, M, C, P : p \cdot Pdefs \rangle$$

### Rule for Procedure calls

$$\langle S, M, (\text{call } P) \cdot C, \dots P : p \dots \rangle \Rightarrow \langle S, M, p \cdot C, \dots P : p \dots \rangle$$

Even though this looks quite innocent, it causes control stack  $C$  *no longer to be of bounded static variation*. To see why, consider say the factorial function defined recursively:

```

Procedure Fac :
  if N = 0 then Result := 1
  else {N := N-1; call Fac; N := N+1; Result := N * Result}

```

The control stack’s depth will be proportional to the initial value of  $N$ , which is dynamic! The result: an unpleasant choice between infinite specialisation on the one hand (for any recursive procedure); or, on the other hand, classifying  $C$  as dynamic, which will result in the existence of source code at run time, and very little speedup.

*Can this problem be solved?* Yes, by using “the trick.” The point is that even though  $Pdefs$  is a dynamic data structure, each of its components comes from the source program  $p$ . This fact can be exploited to modify the interpreter so as gain better specialisation.

A sketch: add another procedure call rule, to make the procedure “return” explicit:

### Rules for Procedure calls

$$\begin{aligned} \langle S, M, (\text{call } P) \cdot C, \dots P : p \dots \rangle &\Rightarrow \langle S, M, p \cdot \text{return} \cdot C, \dots P : p \dots \rangle \\ \langle S, M, \text{return} \cdot C, \dots P : p \dots \rangle &\Rightarrow \langle S, M, C, \dots P : p \dots \rangle \end{aligned}$$

Interestingly, use of “the trick” leads to something close to the target code “return address,” widely used in pragmatic compiler construction. *Claim* (easy to verify): the possible  $C$  contents from **return** down to the next **return** or stack bottom are of  $BSV$ .

A list **Tops** of all such can thus be precomputed from the source program, and so is of  $BSV$ . Now program the **return** transition as follows: pop the top of  $C$  off, and compare it with the elements of **Tops**. Once an equality with **Tops<sub>i</sub>** is found (as it must be,) apply the transitions of Figure 7.1, doing the pattern matching against **Tops<sub>i</sub>**.

The comparison with **Tops** entries will specialise into a series of tests and branches, which realise the procedure return. (Ideally this could be done even faster, by an indirect jump.) In this way all pattern matching will become static, and source text will be used at run time only to implement the return.

The exact form of the source text as used in these comparisons is in fact irrelevant, since its only function is to determine where in the target program control is to be transferred — and so it can be replaced by a *token*, one for each entry in **Tops**. Such tokens are in effect return addresses.

### 7.3 Analysis of the “big-step” interpreter

Interpreter **int-big-step** from Figure 6.2 specialises well. Even though not compositional (as defined in Section 6.2), it is *semicompositional* since the parameters **e**, **ns**, and **pgm** of **Eval** are always substructures of the original source program (not necessarily proper.)

In this case recursion causes no problems; there is no stack, and nothing can grow unboundedly as in the extended instructional interpreter, or as in many small-step semantics.

#### A non-semicompositional variant.

Now consider the interpreter **int-big-step**, extended by a “LET” construction. This is easy to modify by adding an extra CASE:

```
Eval(e, ns, vs, pgm) =   case e of ...

      '(let X = e1 in e2) : Eval(e2, cons('X, ns), cons(v1, vs), pgm)
                           where   v1 = Eval(e1, ns, vs, pgm)
```

With this change, parameter **ns** can take on values that are *not substructures of pgm*. Now **ns** can grow (and without limit) since more deeply nested levels of **let** expressions will give rise to longer lists **ns**. Nonetheless, **ns** *only grows when e shrinks*. Further, it can be reset by a function call, but only to values that are part of **pgm**. Consequently, for any fixed source program, parameter **ns** *cannot*

grow unboundedly as a function of the size of `int-big-step`'s static input `pgm`, and so is of bounded static variation.

## 7.4 Dynamic binding.

Suppose in interpreter `int-big-step` the “function call” construction were implemented differently, as follows:

```
Eval(e, ns, vs, pgm) =    case e of    ...

'(call f es): Eval(e1, append(ns1,ns), append(vs1,vs), pgm) where
    e1 = lookbody(f, pgm)
    ns1 = lookfcns(f, pgm)
    vs1 = Evlist(es, ns, vs, pgm)
```

This can easily be seen to cause `ns` to become of *unbounded static variation*. Intuitively, the reason is that the length of `ns` is no longer tied to any syntactic property of `pgm`. For more technical details, see [5]. Interestingly, the dynamic name binding in the call has an indirect effect on `Eval`. In order to avoid non-termination of specialisation, we must classify the interpreter parameter `ns` as dynamic.

The consequence is (as in Lisp) that target programs obtained by specialising `int-big-step` will contain parameter *names* as well as their values, and function `lookparams` will become a run-time procedure rather than a specialisation-time one, substantially increasing target program running times and space usage.

## 8 On code explosion

Big problems can arise even when everything not classified as dynamic is of bounded static variation!!

### 8.1 Dead static parameters.

A parameter is *semantically dead* at a program point if changes to its value cannot affect the output of the current program. Syntactic approximations to this property are widely used in optimising compilers, especially to minimise register usage.

Now the size of specialised program  $p_s$  critically depends on the number of its *specialised program points*, each of form  $\ell_{(s_1, \dots, s_k)}$  where  $\ell$  is one of  $p$ 's program points, and  $(s_1, \dots, s_k)$  are the values of static parameters whose scope includes point  $\ell$ .

Clearly dead static parameters will merely increase the size of specialised program  $p_s$ , without changing its computation in any way – so removal of dead static parameters can reduce its size. Our first experience with this problem, for a simple imperative language (see [6],) was that not accounting for dead static variables led to a residual program 500 times larger!

## 8.2 Synchronicity in static parameter variation.

**A good example: binary search.** The following imperative program  $p$  performs binary search in an ordered (increasing) table  $T_0, \dots, T_n$  where  $n = 2^m - 1$  and with initial call  $\text{Find}(T, 0, 2^{m-1}, x)$ . Parameter  $x$  is the target to be found,  $i$  points to a position in the table, and  $\text{delta}$  is the size of the portion  $i \dots i + \text{delta}$  of the table currently being compared against.

```
Find(T, i, delta, x) =  
  Loop: if delta = 0 then  
    if x = T[i] then return(i) else return(NOTFOUND);  
    if x >= T[i+delta] then i := i + delta;  
    delta := delta/2;  
    goto Loop
```

We begin by assuming that  $\text{delta}$  is classified as static and that  $i$  is dynamic. Using a term from [6], the program is *weakly oblivious*, since the comparison with  $x$  does not affect the value assigned to  $\text{delta}$ .

Specialising with respect to static initial  $\text{delta} = 2^{3-1} = 4$  and dynamic  $i$  gives program  $p_8$ :

```
if x >= T[i+4] then i := i+4;  
if x >= T[i+2] then i := i+2;  
if x >= T[i+1] then i := i+1;  
if x = T[i] then return(i) else return(NOTFOUND)
```

In general  $p_n$  runs in time  $O(\log(n))$ , and with a better constant coefficient than the original general program. Moreover, it has size  $O(\log(n))$  — acceptable for all but extremely large tables.

**A bad example: binary search.** To illustrate the problems that can occur, consider the same binary search program above with  $n$  static. One may certainly classify both parameters  $\text{delta}$  and  $i$  as static, since both range over  $0, 1, \dots, n-1$ . The resulting program is, however, not oblivious since the test on  $x$  affects the value of static  $i$ . Said differently, parameters  $i$  and  $\text{delta}$  are *independently varying static parameters*.

Specialisation with respect to static initial  $\text{delta} = 4$  and  $i = 0$  now gives the program in Figure 8.1.

The specialised program again runs in time  $O(\log(n))$ , and with a slightly better constant coefficient than above. On the other hand it has size  $O(n)$  — exponentially larger than the weakly oblivious version! This is unacceptably large, except for rather small tables.

## 8.3 Some consequences of binding-time improvements

Binding-time improvements can lead to unexpectedly slow target code. Two typical examples follow.

Interpretation of a program with procedure calls is straightforward if no regard is taken of static-dynamic separation, e.g. a stack containing both source



```

if x >= T[4] then
  if x >= T[6] then
    if x >= T[7] then
      [if x = T[7] then return(7) else return(NOTFOUND)] else
      [if x = T[6] then return(6) else return(NOTFOUND)] else
    if x >= T[5] then
      [if x = T[5] then return(5) else return(NOTFOUND)]
      [if x = T[4] then return(4) else return(NOTFOUND)] else
  if x >= T[2] then
    if x >= T[3] then
      [if x = T[3] then return(3) else return(NOTFOUND)] else
      [if x = T[2] then return(2) else return(NOTFOUND)] else
  if x >= T[1] then
    [if x = T[1] then return(1) else return(NOTFOUND)]
    [if x = T[0] then return(0) else return(NOTFOUND)]

```

Figure 8.1: A large program yielded by specialising binary search.

program text fragments and runtime values may be used. A binding-time improvement can, with some work, split this stack into separated static and dynamic parts. A complication is that the stack depth is dynamic, but “the trick” referred to in [6] can be used to keep the topmost (current) name part of the stack frame static; as sketched in section 7.2.

Target programs produced by specialising such an interpreter may be slower than expected when doing procedure *returns*. The point is that a procedure return is traditionally compiled into an indirect goto. If this is not available, or if the partial evaluator cannot handle it (and few can!), the one-instruction “return jump” will be translated into a series of tests and control transfers, one to each possible return point.

An exactly analogous problem occurs for the lambda calculus: semicompositionality can be ensured by using closures and “the trick.” But the consequence is that every function call will be translated into a series of tests and calls.

#### 8.4 An unfortunate programming style.

The interpreter sketch of Figure 8.2 essentially implements a “big-step” semantics for a tiny imperative language. Nonetheless it does not specialise at all well (though specialisation does in fact terminate.)

This interpreter uses a quite natural idea: if a variable has not already been bound in the store, then it is in effect added, with initial value zero. While it creates no problems for execution, this apparently quite innocent trick can cause catastrophically large target programs to be generated when specialising the interpreter to a source program. The reason is that *mix* must take account of *all execution possibilities* (the “Pachinko syndrome”).

For an example, consider a source program of the following form:

```

Run(pgm, input) = Exec(Cmd1, ns1, cons(input, 'nil), pgm) where
    Cmd1 = lookbody(first_fcn(pgm), pgm)
    ns1  = lookfcns(first_fcn(pgm), pgm)

; Exec : Command x Names x Values x Program -> Names x Values

Exec(Cmd, ns, vs, pgm) = case Cmd of

    '(C1 ; C2)      : Exec(C2, ns, ns1, vs1, pgm) where
        (ns1, vs1) = Exec(C1, ns, vs, pgm),

    '(X := X + 1) : if member(X, ns)
        then (ns, update(X, 1+lookparams(X,ns,vs), ns, vs))
        else (cons(X, ns), cons(1,vs))

    '(IF e THEN C1 ELSE C2) :
        if eval(e0, ns, vs, pgm) then Exec(C1, ns, vs, pgm)
        else Exec(C2, ns, vs, pgm)

Eval(Cmd, ns, vs, pgm) = ...

```

Figure 8.2: Interpreter for a simple imperative language.

```

IF test1 THEN X1 := X1 + 1;
IF test2 THEN X2 := X2 + 1;
...
IF testn THEN Xn := Xn + 1;

```

After `test1` is processed at specialisation time, `mix` will have to allow for either branch to have been taken, so the resulting store = `(ns, vs)` could either have `X1` present or absent. After the second test, 4 possibilities must be accounted for. As a consequence, specialisation will generate a target program whose size is exponential in `n`.

The problem is easily fixed, by revising the interpreter to first scan the entire source program, collecting an initial store in which every variable is bound to zero. The result will yield a target program whose size is proportional to that of the source program. A further optimisation now becomes possible: the type of `Exec` can be simplified to

```

Exec: Command x Names x Values x Program -> Values

```

with correspondingly smaller target code.

## 9 On speedup

Assuming that the specialiser terminates, it is natural to ask: how much speedup can be obtained; and how does the amount of speedup relate to the way the subject program is written?

We begin with a resume (and improvement) of the argument [6] that one can (usually) expect at most speedup by a linear multiplicative factor. We will see, however, that the size of the factor can depend on the static inputs, e.g. one often observes, for instance in pattern matching, that larger static data leads to larger speedups as a function of the size of  $\mathbf{d}$ . The argument below also applies to *deforestation* and *supercompilation*.

### 9.1 Embedding computations into original computations

In most partial evaluators the residual computations can be described more simply than those of the specialiser. Consider a given program  $\mathbf{p}$ , and static and dynamic inputs  $\mathbf{s}$ ,  $\mathbf{d}$ . Let  $\mathbf{p}_s$  be the result of specialising  $\mathbf{p}$  to  $\mathbf{s}$ .

A **first observation** (on all functional or imperative partial evaluators I have seen): there exists an order-preserving *injective mapping*  $\psi$

- from the operations done in the computation of  $\llbracket \mathbf{p}_s \rrbracket(\mathbf{d})$
- to the operations done in the original computation of  $\llbracket \mathbf{p} \rrbracket(\mathbf{s}, \mathbf{d})$ .

The reason is that specialisers normally just sort computations by  $\mathbf{p}$  into *static operations*: those done during specialisation, and *dynamic operations*: those done by the residual program, for which code is generated. Most specialisers do not rearrange the order in which computation is performed, or invent new operations not in  $\mathbf{p}$ , so residual computations can be embedded 1-1 into original ones.

A **second observation, in the opposite direction**: some operations done in the computation of  $\llbracket \mathbf{p} \rrbracket(\mathbf{s}, \mathbf{d})$  can be forced to be residual because they use parts of  $\mathbf{d}$ , which is unavailable at specialisation time. Call such an operation in the  $\llbracket \mathbf{p} \rrbracket(\mathbf{s}, \mathbf{d})$  computation *forced dynamic*. Since it cannot be done by *mix*, residual code must be executed to realise its effect.

It holds for most partial evaluators (but not all) that there exists an order-preserving *surjective order-preserving mapping*  $\phi$  from

- the forced dynamic operations in the computation of  $\llbracket \mathbf{p} \rrbracket(\mathbf{s}, \mathbf{d})$ , onto
- corresponding operations in the computation of  $\llbracket \mathbf{p}_s \rrbracket(\mathbf{d})$ .

**Why most specialisers give at most linear speedup.** Let  $\text{time-force}_p(\mathbf{s}, \mathbf{d})$  be the number of forced dynamic operations done while computing  $\llbracket \mathbf{p} \rrbracket(\mathbf{s}, \mathbf{d})$ . Assuming that all of these *must* be performed by the specialised program, we get a lower bound on its running time:  $\text{time-force}_p(\mathbf{s}, \mathbf{d}) \leq \text{time}_{p_s}(\mathbf{d})$ .

All other  $\mathbf{p}$  operations can be performed without knowing  $\mathbf{d}$ : they depend only on  $\mathbf{s}$ . The maximum length of any sequence of such operations (assuming of course that specialisation terminates) is thus some function  $f(\mathbf{s})$  of the static

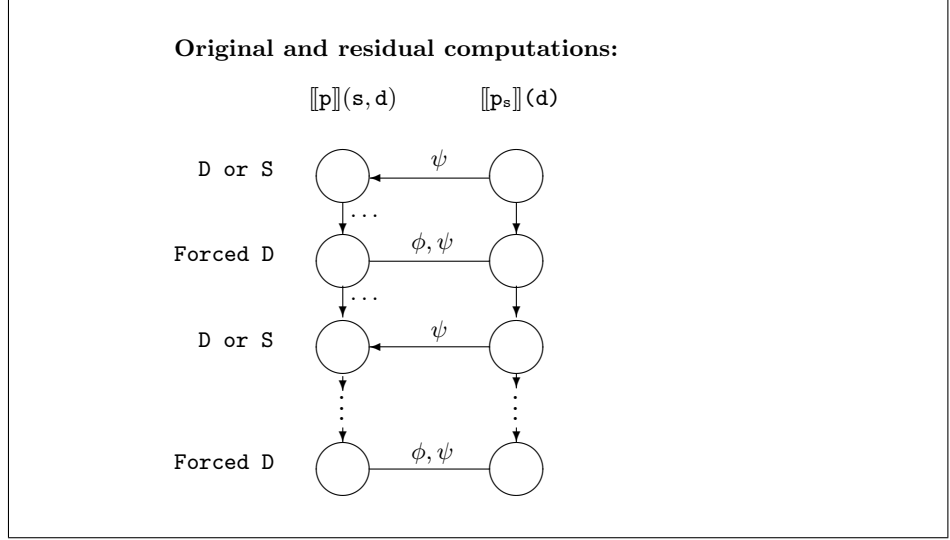


Figure 9.1: Correspondences between original and residual computations.

program input. Thus at most  $f(s)$  static operations occur between any two forced dynamic operations, yielding  $time_p(s, d) \leq (1 + f(s)) \cdot time\_force_p(s, d)$ .

Combining these, as illustrated in Figure 9.1, we see that  $time_p(s, d) \leq (1 + f(s)) \cdot time_{p_s}(d)$ , so

$$speedup_s(d) = \frac{time_p(s, d)}{time_{p_s}(d)} \leq 1 + f(s)$$

This second observation explains the fact that, in general, partial evaluators yield at most linear speedup independent of  $d$ .

## 9.2 Ways to break the linear speedup barrier.

The argument just given rests on the existence of an order-preserving “onto” mapping of forced operations in original computations to residual ones. This is not always the case, and specialisers working differently can obtain superlinear speedup. Three cases follow:

**Dead code elimination:** if `mix` can examine its residual program and discover “semantically dead” code whose execution does not influence the program’s final results, such code can be removed. If it contains dynamic operations, then there will be dynamic  $\llbracket p \rrbracket(s, d)$  operations not corresponding to any  $\llbracket p_s \rrbracket(d)$  operations at all. In this situation, between the two  $p$  operations corresponding to two  $p_s$  operations there can occur a number of  $p$  operations *depending on dynamic input*, which were removed by the dead code elimination technique.

A **logic programming analogy** is the specialiser-time detection of branches in the Prolog’s *SLD computation tree* which are guaranteed to fail (and may

depend on dynamic data.) This is much more significant than in functional programming, in which time-consuming sequences of useless operations are most likely a sign of bad programming — whereas the ability to prune useless branches of search trees is part of the essence of logic programming.

**Removing repeated subcomputations** is another way to achieve super-linear speedup, and somehow seems less trivial than just eliding useless computations. The program transformation literature abounds with large speedups achieved by using memoisation to avoid recomputation of results, the standard one being to go from the exponential-time Fibonacci function to a linear-time (or even logarithmic-time) one. Integrating such more powerful transformations into fully automated partial evaluators remains a challenging problem, and in my opinion one well worth attempting.

## 10 Summary: what to do and not to do ...

By hook or by crook, ensure that *every* static parameter is of BSV. To do this, recall from Section 4.1 the “Pachinko principle:” that specialisation must account for *all possible computation paths* on any one set of static inputs, given *all possible dynamic inputs* coupled with these inputs. More concretely:

1. Write your interpreter compositionally if possible, or at least semicompositionally (Section 5.)
2. If practical, write it by transliterating a big-step operational (Section 6.3) or a denotational semantics (Section 6.2.) This usually gives at least semicompositionality.
3. If your interpreter is based on a runtime execution model, you will have to do your own analyses concerning which parameters are of BSV. Beware of data structures that contain static data but can *grow unboundedly under dynamic control*. An example: the control stack *C* of Section 7.2.
4. If your interpreter is based on a small-step operational semantics (Section 6.3), first ensure *determinism* and *uniqueness of redexes*.  
(Here I’m assuming the implementation language is deterministic; if your specialiser handles nondeterministic languages, it is ahead of the state of the art as I know it.)  
Once determinism and redex uniqueness are established, ensure that the set of evaluation contexts is finite for any one source program. One way is by introducing closures or similar devices to ensure “dual semicompositionality.”
5. In order to avoid *code explosion*, ensure that at any program point:
  - There are no *dead static variables* (Section 8.1.)
  - There are no (or as few as possible) *independently varying static parameters* (Section 8.2.)
6. Don’t expect *superlinear speedup* (as a function of the dynamic input size), unless the specialiser used is rather sophisticated (Section 9.1.)
7. If you wish to achieve superlinear speedup on an insufficiently sophisticated specialiser, then *write your interpreter* to incorporate techniques statically

recognisable optimisation such as those seen in Section 9.2 (memoisation, static detection of semantically dead code, etc.)

**Acknowledgements** Warm thanks are due to many in the Topps group at DIKU, especially Arne Glenstrup; and to participants in the 1996 Partial Evaluation Dagstuhl meeting, in particular Peter Thiemann, Olivier Danvy, and Robert Glück.

## References

1. A. Bondorf and O. Danvy, ‘Automatic autoprojection of recursive equations with global variables and abstract data types,’ *Science of Computer Programming*, 16:151–195, 1991.
2. C. Consel, ‘New insights into partial evaluation: The Schism experiment,’ in H. Ganzinger (ed.), *ESOP ’88, 2nd European Symposium on Programming, Nancy, France, March 1988 (Lecture Notes in Computer Science, vol. 300)*, pp. 236–246, Berlin: Springer-Verlag, 1988.
3. A. De Niel, E. Bevers, and K. De Vlaminck, ‘Partial evaluation of polymorphically typed functional languages: The representation problem,’ in M. Billaud et al. (eds.), *Analyse Statique en Programmation Équationnelle, Fonctionnelle, et Logique, Bordeaux, France, Octobre 1991 (Bigre, vol. 74)*, pp. 90–97, Rennes: IRISA, 1991.
4. J. Gallagher: Specialization of logic programs. PEPM 93 (Partial Evaluation and Semantics-based Program Manipulation), pp. 88–98. ACM Press, 1993.
5. N.D. Jones, A.G. Glenstrup ‘BTA Algorithms to ensure termination of offline partial evaluation,’ submitted to *Second Andrei Ershov Memorial Conference, Akademgorodok, Russia*, June 1996.
6. Neil D. Jones, C.K. Gomard, P. Sestoft, *Partial Evaluation and Automatic Program Generation*, Prentice Hall International Series in Computer Science, 1993.
7. J.W. Lloyd, J.C. Shepherdson, ‘Partial evaluation in logic programming,’ *Journal of Logic Programming* 11(3-4), pp. 217–242, 1991.
8. T. Mogensen, ‘Partially static structures in a self-applicable partial evaluator,’ in D. Bjørner, A.P. Ershov, and N.D. Jones (eds.), *Partial Evaluation and Mixed Computation*, pp. 325–347, Amsterdam: North-Holland, 1988.
9. T. Mogensen, ‘Separating binding times in language specifications,’ in *Fourth International Conference on Functional Programming Languages and Computer Architecture, London, England, September 1989*, pp. 14–25, Reading, MA: Addison-Wesley, 1989.
10. D. Sahlin, ‘The Mixtus approach to automatic partial evaluation of full Prolog,’ in S. Debray and M. Hermenegildo (eds.), *Logic Programming: Proceedings of the 1990 North American Conference, Austin, Texas, October 1990*, pp. 377–398, Cambridge, MA: MIT Press, 1990.
11. D.A. Schmidt, *Denotational Semantics*, Boston, MA: Allyn and Bacon, 1986.
12. V.F. Turchin, ‘The concept of a supercompiler,’ *ACM Transactions on Programming Languages and Systems*, 8(3):292–325, July 1986.
13. D. Weise, R. Conybeare, E. Ruf, and S. Seligman, ‘Automatic online partial evaluation,’ in J. Hughes (ed.), *Functional Programming Languages and Computer*

*Architecture, Cambridge, Massachusetts, August 1991 (Lecture Notes in Computer Science, vol. 523)*, pp. 165–191, Berlin: Springer-Verlag, 1991.