

A Roadmap to Metacomputation by Supercompilation

Robert Glück & Morten Heine Sørensen

DIKU, Department of Computer Science, University of Copenhagen
Universitetsparken 1, DK-2100 Copenhagen Ø, Denmark
E-mail: {glueck,rambo}@diku.dk

Dedicated to V.F. Turchin on the Occasion of his 65th Birthday

Abstract. This paper gives a gentle introduction to Turchin’s supercompilation and its applications in metacomputation with an emphasis on recent developments. First, a complete supercompiler, including positive driving and generalization, is defined for a functional language and illustrated with examples. Then a taxonomy of related transformers is given and compared to the supercompiler. Finally, we put supercompilation into the larger perspective of metacomputation and consider three metacomputation tasks: specialization, composition, and inversion.

Keywords: Program transformation, supercompilation, driving, generalization, metacomputation, metasystem transition.

1 Introduction

Over the years a number of automatic program transformers from Burstall and Darlington’s unfold/fold framework [11] have been devised and implemented. The most popular is *partial evaluation* which performs *program specialization*. The possibility, in principle, of partial evaluation is contained in Kleene’s *s-m-n* Theorem [47]. The idea to use partial evaluation as a *programming tool* can be traced back to work beginning in the late 1960’s by Lombardi and Raphael [58, 57], Dixon [18], and Chang and Lee [12]. Important contributions were made in the seventies by Futamura [23, 24], by Sandewall’s group [6], by Ershov [19, 20], and later by Jones’ group [45, 46]. In the eighties program specialization became a research field of its own, *e.g.* [8, 14, 16, 44, 53].

Supercompilation [87], conceived by Turchin in the early seventies in Russia for the programming language Refal, achieves the effects of partial evaluation as well as more dramatic optimizations. Turchin formulated the transformations necessary for supercompilation, including the central *rule of driving* and the *outside-in strategy*, in 1972 [80, 81] and the main results concerning self-application, *metasystem transition*, in 1973. The book [96] defined all three *Futamura projections* in terms of metasystem transition. In the English language, the work on supercompilation was first described in [83, 84, 85, 86] and then developed further in [87, 89, 93, 99]. For more historical details, see Turchin’s personal account [95]. Despite these remarkable contributions, supercompilation has not found recognition outside a small circle of experts.

This paper gives a gentle introduction to the principles of supercompilation in terms of a *positive supercompiler* [34, 73, 74, 75, 76] comprising two components, *driving* (Sect. 2) and *generalization* (Sect. 3). The supercompiler is compared to related program transformers (Sect. 4), and put into the larger perspective of metacomputation (Sect. 5). We give references to the literature throughout the text, which can hopefully be used as a starting point for further reading. The bibliography contains a comprehensive list of Russian and English titles.

1.1 Object Language

We are concerned with a first-order functional language; the intended operational semantics is normal-order graph reduction to weak head normal form [7].

The syntax of our language appears in Fig. 1 (where $m > 0, n \geq 0$). We assume denumerable, disjoint sets of symbols for variables $v \in \mathcal{V}$, constructors $c \in \mathcal{C}$, and functions $f \in \mathcal{F}$ and $g \in \mathcal{G}$; symbols all have fixed arity. A given program makes use of a finite number of different symbols.

$\mathcal{Q} \ni q$	$::= d_1 \dots d_m$	(program)
$\mathcal{D} \ni d$	$::= f v_1 \dots v_n \triangleq t$	(f-function definition, no patterns)
	$\quad \quad g p_1 v_1 \dots v_n \triangleq t_1$	
	$\quad \quad \quad \vdots$	(g-function definition with patterns)
	$\quad \quad \quad g p_m v_1 \dots v_n \triangleq t_m$	
$\mathcal{T} \ni t$	$::= v$	(variable)
	$\quad \quad c t_1 \dots t_n$	(constructor)
	$\quad \quad f t_1 \dots t_n$	(f-function call)
	$\quad \quad g t_0 t_1 \dots t_n$	(g-function call)
	$\quad \quad \text{if } t_1 = t_2 \text{ then } t_3 \text{ else } t_4$	(conditional with equality test)
$\mathcal{P} \ni p$	$::= c v_1 \dots v_n$	(flat pattern)

Fig. 1. Syntax of programs, definitions, terms, and patterns.

A program $q \in \mathcal{Q}$ is a sequence of function definitions $d \in \mathcal{D}$ where the right side of each definition is a term $t \in \mathcal{T}$ constructed from variables, constructors, function calls, and conditionals. We require that no two patterns p_i and p_j in a g-function definition contain the same constructor c , that no variable occur more than once in a pattern, and that all variables on the right side of a definition be present in its left side. Figure 2 shows the function a for appending two lists, using the short notation $[]$ and $(x : xs)$ for the list constructors nil and $cons\ x\ xs$.

$a []\ ys$	$\triangleq ys$
$a (x : xs)\ ys$	$\triangleq x : a\ xs\ ys$

Fig. 2. Example program *append*.

Our language contained case-expressions in [76, 74], because the connection between positive supercompilation and deforestation is clearest for case-expressions. However, g -functions [21] lead to a simpler presentation of generalization.

2 Driving

Driving takes a term and a program and constructs a possibly infinite *process tree*, representing all possible computations with the term, in a certain sense. Figure 3 shows part of the infinite process tree for the term $a (a xs ys) xs$ (note the repeated variable xs). Each node contains a term t and its children contain

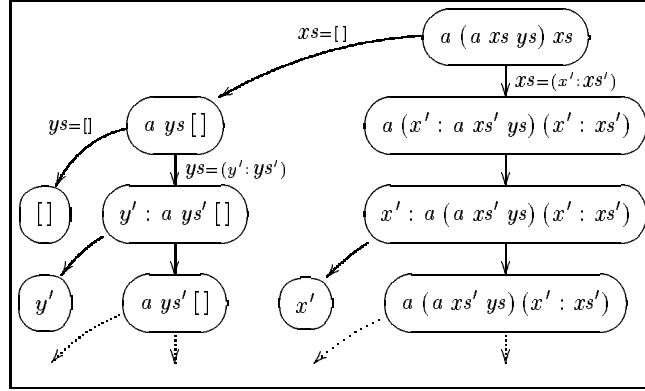


Fig. 3. Example process tree for $a (a xs ys) xs$.

the terms that arise by one *normal-order reduction* step from t . Whenever the reduction step has different possible outcomes there are several children so as to account for all possibilities. For instance, the topmost branching in Fig. 3 corresponds to the cases $xs = []$ and $xs = (x' : xs')$. Note how the terms in the rightmost branches continue to grow due to the unification-based information propagation of driving, *e.g.* xs is replaced by $(x' : xs')$.

In Sect. 2.1 we define normal-order reduction, and in Sect. 2.2 we introduce process trees and define driving.

2.1 Normal-Order Reduction

A *value* is a term built exclusively from constructors and variables. An *observable* is either a variable or a term with a known outermost constructor. Any term which is not an observable can be decomposed into the form $e[r]$ where the *redex* r is the outermost reducible subterm and the *evaluation context* e is the surrounding part of the term.

More precisely, define *values*, *observables*, *redexes*, and *evaluation contexts* by the syntactic classes \mathcal{B} , \mathcal{O} , \mathcal{R} and \mathcal{E} , respectively, as in Fig. 4. Define $e[t]$ to be the result of substituting t for the “hole” \diamond in e .

$\mathcal{B} \ni b$	$::= v \mid c b_1 \dots b_n$	(value)
$\mathcal{O} \ni o$	$::= v \mid c t_1 \dots t_n$	(observable)
$\mathcal{R} \ni r$	$::= f t_1 \dots t_n$ $\mid g o t_1 \dots t_n$ $\mid \text{if } b_1 = b_2 \text{ then } t_1 \text{ else } t_2$	(redex)
$\mathcal{E} \ni e$	$::= \diamond$ $\mid g e t_1 \dots t_n$ $\mid \text{if } d = t_2 \text{ then } t_3 \text{ else } t_4$ $\mid \text{if } b = d \text{ then } t_3 \text{ else } t_4$	(evaluation context)
d	$::= e \mid c b_1 \dots b_{i-1} d t_{i+1} \dots t_n$	

Fig. 4. Values, observables, redexes, evaluation contexts.

Lemma 1 (the unique decomposition property). For any $t \in \mathcal{T}$ either $t \in \mathcal{O}$ or there exists a unique pair $(e, r) \in \mathcal{E} \times \mathcal{R}$ such that $t \equiv e[r]$.

Figure 5 shows example decompositions. In (1) the outermost call to the f -function f can be unfolded; the evaluation context is empty. In (2) the call to f has to be unfolded (the call to the g -function g cannot be unfolded because the term $f t$ does not have a known outermost constructor). In (3) the call to f has to be unfolded; both sides of an equality test must be values.

	t	e	r
(1)	$f (g t)$	\diamond	$f (g t)$
(2)	$g (f t)$	$g \diamond$	$f t$
(3)	$\text{if } x = c (f t) \text{ then } t' \text{ else } t''$	$\text{if } x = c \diamond \text{ then } t' \text{ else } t''$	$f t$

Fig. 5. Examples of decomposition into redex and evaluation context.

The rules for *normal-order reduction* are given by the map \mathcal{N} from terms to ordered sequences $\langle t_1, \dots, t_n \rangle$ of terms in Fig. 6. The rules of \mathcal{N} are mutually exclusive and together exhaustive by the unique decomposition property.

Notation: the expression $t\{v_i := t_i\}_{i=1}^n$ denotes the result of simultaneously replacing all occurrences of v_i by the corresponding term t_i where $1 \leq i \leq n$. The expression $[b, b']$ denotes an idempotent most general unifier $\{v_i := t_i\}_{i=1}^n$ of b and b' if it exists, and *fail* otherwise, where we stipulate $t \text{ fail} \equiv t$. A value is *ground* if it contains no variables. To avoid name capture, the variables occurring in definitions, like $f v_1 \dots v_n \triangleq s$ in the third clause, must be fresh.

Note the substitutions in the two last clauses. The assumed outcome of the test is propagated to the terms resulting from the step. We call this *unification-based* information propagation (*cf.* Sect. 4).

2.2 Process Trees

A *process tree* is a tree where each node is labeled with a term t and all edges leaving a node are ordered. Every node may have an additional *mark*.

t	$\mathcal{N}[\![t]\!]$
x	$\langle \rangle$
$c\ t_1 \dots t_n$	$\langle t_1, \dots, t_n \rangle$
$e[f\ t_1 \dots t_n]$	$\langle e[s\{v_i:=t_i\}_{i=1}^n] \rangle$ if $f\ v_1 \dots v_n \triangleq s$
$e[g\ (c\ t_1 \dots t_i)\ t_{i+1} \dots t_n]$	$\langle e[s\{v_i:=t_i\}_{i=1}^n] \rangle$ if $g\ (c\ v_1 \dots v_i)\ v_{i+1} \dots v_n \triangleq s$
$e[g\ x\ t_1 \dots t_n]$	$\langle (e[s_1\{v_i:=t_i\}_{i=1}^n])\{x:=p_1\}, \dots, (e[s_m\{v_i:=t_i\}_{i=1}^n])\{x:=p_m\} \rangle$ if $g\ p_1\ v_1 \dots v_n \triangleq s_1; \dots ; g\ p_m\ v_1 \dots v_n \triangleq s_m$
$e[\text{if } b=b' \text{ then } t \text{ else } t']$	$\begin{cases} \langle e[t] \rangle & \text{if } b, b' \text{ are ground, } b \equiv b' \\ \langle e[t'] \rangle & \text{if } b, b' \text{ are ground, } b \not\equiv b' \\ \langle (e[t])[b, b'], e[t'] \rangle & \text{if } b, b' \text{ are not both ground} \end{cases}$

Fig. 6. Normal-order reduction step.

Definition 2. Let T be a process tree and \textcircled{t} an unmarked leaf node in T . Then $UNFOLD(T, \textcircled{t})$ is the process tree obtained by marking \textcircled{t} and adding n unmarked children labeled t_1, \dots, t_n , where $\mathcal{N}[\![t]\!]=\langle t_1, \dots, t_n \rangle$.

Driving is the action of constructing process trees using two essential principles: normal-order strategy and unification-based information propagation.

Algorithm 3 (driving.)

1. INPUT $t_0 \in \mathcal{T}$, $q \in \mathcal{Q}$
2. LET T_0 be the process tree with unmarked node labeled t_0 . SET $i = 0$.
3. WHILE there exists an unmarked leaf node N in T_i :
 - (a) $T_{i+1} = UNFOLD(T_i, N)$
 - (b) SET $i = i + 1$
4. OUTPUT T_i

3 A Positive Supercompiler

In the previous section we used driving to construct a potentially infinite process tree. The purpose of *generalization* is to ensure that one constructs instead a finite *partial process tree* from which a new term and program can be recovered.

The idea is that if a leaf node M has an ancestor L and it “seems likely” that continued driving will generate an infinite sequence L, \dots, M, \dots then M should not be driven any further; instead we should perform *generalization*. In Sect. 3.1 we define a criterion, a so-called *whistle*, that formalizes the decision when to stop. In Sect. 3.2 we introduce some notions that are used in Sect. 3.3 to define generalization. This culminates in a definition of a *positive supercompiler*.

3.1 When to Stop?

We stop driving at a leaf node with label t if one of its ancestors has label s and $s \sqsubseteq t$, where \sqsubseteq is the *homeomorphic embedding relation* known from term algebra [17]. Variants of this relation are used in termination proofs for term rewrite systems [17] and for ensuring local termination of partial deduction [9]. After it was taken up in [76], it has inspired more recent work [5, 55, 94, 95].

The rationale behind this relation is that in any infinite sequence t_0, t_1, \dots that arises during driving of a program, there *definitely* exists some $i < j$ with $t_i \sqsubseteq t_j$, so driving cannot proceed infinitely. Moreover, if $t_i \sqsubseteq t_j$ then all the subterms of t_i are present in t_j embedded in extra subterms. This suggests that t_j might arise from t_i by some infinitely continuing system, so driving will be stopped for a good reason.

The *homeomorphic embedding* \sqsubseteq is the smallest relation on \mathcal{T} satisfying the rules in Fig. 7, where $h \in \mathcal{X} \cup \mathcal{C} \cup \mathcal{F} \cup \mathcal{G} \cup \{\text{ifthenelse}\}$, $x, y \in \mathcal{X}$, and $s, s_i, t \in \mathcal{T}$.

Variable	Diving	Coupling
$x \sqsubseteq y$	$\frac{s \sqsubseteq t_i \text{ for some } i}{s \sqsubseteq h(t_1, \dots, t_n)}$	$\frac{s_1 \sqsubseteq t_1, \dots, s_n \sqsubseteq t_n}{h(s_1, \dots, s_n) \sqsubseteq h(t_1, \dots, t_n)}$

Fig. 7. Homeomorphic embedding.

Diving detects a subterm embedded in a larger term, and *coupling* matches the subterms of two terms. Some examples and non-examples appear in Fig. 8. It is not hard to give an algorithm $WHISTLE(M, N)$ deciding whether $M \sqsubseteq N$.

$b \sqsubseteq a(b)$	$a(c(b)) \not\sqsubseteq c(b)$
$c(b) \sqsubseteq c(a(b))$	$a(c(b)) \not\sqsubseteq c(a(b))$
$d(b, b) \sqsubseteq d(a(b), a(b))$	$a(c(b)) \not\sqsubseteq a(a(a(b)))$

Fig. 8. Examples and non-examples of embedding.

3.2 Most Specific Generalization

We define the generalization of two terms t_1, t_2 as the most specific generalization (msg) $[t_1, t_2]$. A well-known result in term algebra states that any two $t, s \in \mathcal{T}$ have an msg which is unique up to renaming. Examples are shown in Fig. 9.

s t	t_g	θ_1	θ_2
$b \sqsubseteq a(b)$	x	$\{x := b\}$	$\{x := a(b)\}$
$c(b) \sqsubseteq c(a(b))$	$c(x)$	$\{x := b\}$	$\{x := a(b)\}$
$c(y) \sqsubseteq c(a(y))$	$c(x)$	$\{x := y\}$	$\{x := a(y)\}$
$d(b, b) \sqsubseteq d(a(b), a(b))$	$d(x, x)$	$\{x := a(b)\}$	$\{x := a(b)\}$

Fig. 9. Examples of most specific generalization.

Definition 4 (instance, generalization, msg, distinct). Given $t_1, t_2 \in \mathcal{T}$.

1. An *instance* of t_1 is a term of the form $t_1\theta$ where θ is a substitution.
2. A *generalization* of t_1, t_2 is a triple $(t_g, \theta_1, \theta_2)$ where $t_g\theta_1 \equiv t_1$ and $t_g\theta_2 \equiv t_2$.
3. A generalization $(t_g, \theta_1, \theta_2)$ of t_1 and t_2 is *most specific* (msg) if for every generalization $(t'_g, \theta'_1, \theta'_2)$ of t_1 and t_2 it holds that t_g is an instance of t'_g .
4. Two terms t_1 and t_2 are *disjoint* if their msg is of the form (x, θ_1, θ_2) .

Algorithm 5 (msg.) An msg $[s, t]$ of $s, t \in \mathcal{T}$ is computed by exhaustively applying the rewrite rules in Fig. 10 to the initial triple $(x, \{x := s\}, \{x := t\})$:

$$\boxed{\begin{array}{l} \left(\begin{array}{l} t_g \\ \{x := h(s_1, \dots, s_n)\} \cup \theta_1 \\ \{x := h(t_1, \dots, t_n)\} \cup \theta_2 \end{array} \right) \rightarrow \left(\begin{array}{l} t_g\{x := h(y_1, \dots, y_n)\} \\ \{y_1 := s_1, \dots, y_n := s_n\} \cup \theta_1 \\ \{y_1 := t_1, \dots, y_n := t_n\} \cup \theta_2 \end{array} \right) \\ \\ \left(\begin{array}{l} t_g \\ \{x := s, y := s\} \cup \theta_1 \\ \{x := t, y := t\} \cup \theta_2 \end{array} \right) \rightarrow \left(\begin{array}{l} t_g\{x := y\} \\ \{y := s\} \cup \theta_1 \\ \{y := t\} \cup \theta_2 \end{array} \right) \end{array}}$$

Fig. 10. Computing most specific generalizations.

3.3 Partial Process Trees

A *partial process tree* differs from a process tree in that it may contain an extra kind of nodes, *generalization-nodes*, with label of form **let** $x_1=t_1 \dots x_n=t_n$ **in** t and $n+1$ children labeled t_1, \dots, t_n, t , respectively, where x_1, \dots, x_n do not occur in t_1, \dots, t_n . This kind of node has the distinct feature that the $n+1$ 'st edge may go to an ancestor of the node instead of going to a child; such an edge is called a *return edge*. We regard a partial process tree as an acyclic graph by ignoring return edges, so *ancestor*, *leaf*, etc. apply only to non-return edges. The labels on generalization nodes are unrelated to all other labels wrt. \leq .

The operations in the following definition, inspired by [59], appear in Fig. 11.

Definition 6. Let T be a partial process tree with node (t) with ancestor (s) .

1. If t is an instance of s , i.e. $t \equiv s\{x_1:=t_1, \dots, x_n:=t_n\}$, then $FOLD(T, (s), (t))$ is the tree obtained as follows. Replace (t) by $(\text{let } x_1=t_1 \dots x_n=t_n \text{ in } s)$ which is marked, has return edge to (s) , and n unmarked children $(t_1), \dots, (t_n)$.
2. If $[s, t] = (t_g, \{x_1:=t_1, \dots, x_n:=t_n\}, \theta)$, then $GENERALIZE(T, (s), (t))$ is the partial process tree obtained as follows. Delete all descendants of (s) , and replace (s) by $(\text{let } x_1=t_1 \dots x_n=t_n \text{ in } t_g)$ with a mark and $n+1$ unmarked children $(t_1), \dots, (t_n), (t_g)$. Return edges from (s) or its descendants are erased.
3. If $t \equiv ht_1 \dots t_n$ then $SPLIT(T, (s), (t))$ is the partial process tree obtained as follows. Let $t' \equiv h\ x_1, \dots, x_n$ where x_1, \dots, x_n are new variables, replace (t) by $(\text{let } x_1=t_1 \dots x_n=t_n \text{ in } t')$ which has a mark and $n+1$ unmarked children $(t_1), \dots, (t_n), (t')$.




Fig. 11. Operations on partial process tress.

Algorithm 7 (positive supercompilation.)

1. INPUT $t_0 \in \mathcal{T}$, $q \in \mathcal{Q}$
2. LET T_0 be the partial process tree with unmarked node labeled t_0 . SET $i := 0$.
3. WHILE there exists an unmarked leaf node N in T_i :
 - (a) IF there exists no ancestor M such that $WHISTLE(M, N)$
 THEN $T_{i+1} := UNFOLD(T_i, N)$
 ELSE
 - i. LET M be an ancestor such that $WHISTLE(M, N)$
 - ii. IF node N is an instance of M THEN $T_{i+1} := FOLD(T_i, M, N)$
 ELSE IF N and M are disjoint THEN $T_{i+1} := SPLIT(T_i, M, N)$
 ELSE $T_{i+1} := GENERALIZE(T_i, M, N)$
 - (b) SET $i := i + 1$
4. OUTPUT T_i

Theorem 8. *Algorithm 7 always terminates.*

The theorem comprises two important facts. The first, a consequence of Kruskal's Tree Theorem, see [17], is that driving a branch cannot go on indefinitely without interruption by generalization. The second fact is that whenever generalization happens, the new nodes are all smaller in a certain order (involving term size and number of occurrences of the same variable) than the one they replace. This would not be the case if the *GENERALIZE* operation were performed when the two compared nodes have labels that are disjoint terms.

As for correctness, it is easily proved that each step of the transformation rules preserves normal-order graph reduction semantics; extending rigorously the proof to account for folding is more involved. A general technique due to Sands [71] can be used to prove this for (positive) supercompilation, see [70].

3.4 Discussion of the Algorithm

A number of choices are left open or settled in an arbitrary way in our algorithm.

First, our algorithm follows Turchin's *generalization principle* [89] which states that a generalization between two terms has a meaning only in the context of the computation process in which they take part. Indeed, our algorithm searches only the ancestors of a leaf node. However, to avoid the generation of duplicate definitions one might imagine searching across different branches; see *e.g.* [38].

Second, our algorithm does not specify a particular strategy for *selecting unmarked leaf nodes*. One may chose a breath-first or depth-first strategy.

Third, in case driving stops the algorithm may employ different strategies for *selecting ancestors* for generalization. For instance, one may choose the closest ancestor, or the ancestor that gives the most specific generalization.

Fourth, when we perform a $GENERALIZE(T, M, N)$ step we replace node M . Instead one could replace N , since this avoids destroying the whole subtree with root M ; other branches from M can be retained with no loss of information.

Fifth, the operator $[\bullet, \bullet]$ and the stop criterion can be varied (as in [55]). The operation $SPLIT(T, s, t)$ may be refined to split t in another way; *e.g.* if $s \equiv h\ x$

and $t \equiv l(k(h y))$ then split such that $t' \equiv l(k z)$ (Turchin's algorithm [89] maintains a stack structure of common contexts to determine split points).

Finally, one can imagine various optimizations of which we will discuss only one, namely *transient reductions*. A term of the form $e[g x t_1 \dots t_n]$ is *non-transient*, all other terms are *transient*. The optimization consists in adding the disjunction “or the label of N is transient” to the condition in (3a) of Algorithm 7. So only terms that require a choice at run-time are compared to ancestors for whistling in the partial process tree. The rationale is that any loop in the program must pass through a choice point unless there is an unconditional loop in the program. However, this means that the partial process tree in principle can be infinite—a risk considered worth taking in the area of partial evaluation [44]. An alternative to transient reductions that gives similar residual programs, and at the same time retains the guarantee that the constructed partial process trees are finite, are *characteristic trees*, e.g. [55].

The partial process tree for $a(a xs ys) xs$ using transient reductions appears in Fig. 12.

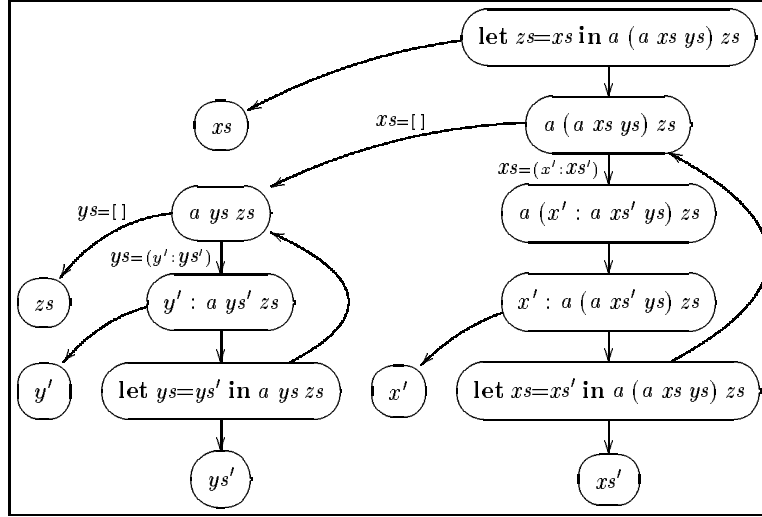


Fig. 12. Example partial process tree with transient reductions.

From this tree one can generate the term $a' xs ys xs$ and a new program (Fig. 13). This transformation is noteworthy because the initial term requires passing the list xs twice, whereas the new term passes xs only once.

$a' [] ys zs$	\triangleq	$a'' ys zs$
$a' (x : xs) ys zs$	\triangleq	$x : a' xs ys zs$
$a'' [] zs$	\triangleq	zs
$a'' (y : ys) zs$	\triangleq	$y : a'' ys zs$

Fig. 13. More efficient double append program.

3.5 Comparative Remarks

From its very inception, supercompilation has been tied to a specific programming language, called *Refal* [78], a language inspired by Markov algorithms. A Refal program is a sequence of rewrite rules, used to transform data in the form of associative and possibly nested symbol strings and offers certain advantages for programming, *e.g.* [90, 62]. Running interpreters were available by the end of the 1960's [22]; different versions of the language were implemented [96, 50, 90, 40].

Driving and generalization for our language are simplified considerably due to simpler data structures, untyped variables, and flat patterns (essentially *elementary contractions* [89]). Due to Refal's data structure most general unifiers do not always exist; a *generalized matching algorithm* is defined in [84, 87]. An optimization technique for pattern matching in Refal was developed in [51].

We should note that supercompilation is a normal-order transformation that can be applied to programs with call-by-value semantics, and that transformed programs can be interpreted call-by-value, as done by Turchin. As a result, supercompilation may make programs terminate more often. Here supercompilation is used to transform programs with normal-order graph reduction semantics into programs with the same semantics and the same termination properties.

Process trees correspond to Turchin's *graph of states* [85] (sometimes called Refal graphs or pattern matching graphs [95]). Driving was used in the seventies in a system for inverse computation, called URA (*cf.* [65]; Sect. 5.2), an application that has been suggested earlier in [80]. *Neighborhood analysis* [84] uses driving to determine sets of data that pass through a computation process in identical ways; suggested for generalization [89] and program testing [2, 3]. Several supercompilers have been developed for Refal [39, 49, 67, 89, 98, 99] including several experimental systems by the Refal group in Moscow (mostly unpublished, except [4]). The first 'non-Refal' supercompiler was [34].

4 Related Program Transformers

In this section we compare positive supercompilation briefly to *partial evaluation*, *deforestation*, *partial deduction*, *perfect supercompilation*, and *generalized partial computation*. First we introduce a number of axes along which transformers can be compared, and then enter the coordinates of the above transformers.

4.1 Some Dimensions in Automatic Program Transformation

Information propagation. Every program transformer maintains a certain level of information propagation; we consider *constant propagation*, *unification-based information propagation*, and *constraint-based information propagation*. The three levels differ in how much information is recorded about pattern matching and tests, corresponding to the transformation rules in Fig. 14.

In constant propagation the outcome of tests are ignored. In unification-based propagation substitutions into the transformed terms are used to represent

	$T\llbracket \text{if } u=v \text{ then } t \text{ else } s \rrbracket =$	<i>information propagation</i>
(a)	$\text{if } u=v \text{ then } T\llbracket t \rrbracket \text{ else } T\llbracket s \rrbracket$	constant propagation
(b)	$\text{if } u=v \text{ then } T\llbracket t\{u:=v\} \rrbracket \text{ else } T\llbracket s \rrbracket$	unification-based
(c)	$\text{if } u=v \text{ then } T\llbracket t \rrbracket\{u=v\} \text{ else } T\llbracket s \rrbracket\{u \neq v\}$	constraint-based

Fig. 14. Information propagation.

the outcome of tests. In constraint-based propagation the transformer explicitly maintains sets of constraints recording previous tests (*restrictions* [84, 34]). Depending on the programming language other abstract properties may be propagated, *e.g.* [72, 15, 77, 43].

Evaluation strategy. One can view a program transformer as an extension of an interpreter, *e.g.* [34, 31, 61, 74]. This implies that the transformer has an *evaluation strategy* that it inherits from the underlying interpreter. More concretely, the transformer processes nested function calls in some order. We consider transformers that use *inside-out* (or *call-by-value* or *applicative order*) and *outside-in* (or *call-by-name* or *normal-order*).

Control restructuring. Control restructuring is concerned with the relationship between program points in the subject and the residual program [10, 69]:

Monovariant: any program point in the subject program gives rise to zero or one program point in the residual program.

Polyvariant: any program point in the subject program can give rise to one or more program points in the residual program.

Monogenetic: any program point in the residual program is produced from a single program point of the subject program.

Polygenetic: any program point in the residual program may be produced from one or more program points of the subject program.

4.2 A Taxonomy of Transformers

Deforestation, due to Wadler [100], performs *program composition* by eliminating intermediate data structures. Deforestation performs, as a special case, program specialization (see [75]). Deforestation is very similar to positive supercompilation except that it uses constant propagation rather than unification-based information propagation. Also, it does not incorporate generalization during transformation. In the original version, deforestation is guaranteed to terminate for a certain class of programs without generalization. Later extensions use annotations computed by a static analysis to guide generalization.

Partial evaluation performs *program specialization* and, as presented in [44], uses only constant propagation [34, 75, 74]. This limitation applies to all variants of partial evaluation: offline and online approaches with and without partially static structures. The usual evaluation strategy for partial evaluators is applicative-order, see [61].

Partial deduction, as in [56, 52, 26], and positive supercompilation have essential aspects in common [38]: the way in which goals are unified and how the

resulting substitutions are applied to the goals in the next transformation step (construction of a partial SLDNF tree), is much like in the clauses of driving.

Since in logic programs predicates cannot occur inside predicates, there is no direct correspondence to the rules for nested function calls which achieve deforestation. However, local variables in logic programs often represent intermediate data structures that could be removed by more sophisticated techniques. Partial deduction in logic programming is not capable of removing them; this requires an extension of the techniques, see *e.g.* [63, 54, 33].

Turchin's supercompiler [87] and our positive supercompiler are identical with respect to the propagation of positive information, except for certain trivial differences. The main difference between the two is that the former also maintains *negative information*, *i.e.* the information that a test failed, and this is maintained in the form of constraints (see *perfect driving* [34]).

Generalized partial computation (GPC), due to Futamura [25], has a similar effect and power as supercompilation, but has arbitrary tests rather than just patterns and equality tests. The underlying logic for the tests can be any logic system, for example predicate logic, and may be undecidable for certain logic formulas. In this view, positive supercompilation can be seen as propagating structural predicates that can always be resolved. Abstract interpretation can be used to propagate additional information, *e.g.* [43, 64].

These observations are summarized in Fig. 15. For a more detailed discussion on information propagation see [34, 75, 38, 74, 76], and for more on evaluation strategies see [13, 61]. These papers also give examples of optimizations that require the transformer to use a specific evaluation strategy or level of information propagation. For instance, to pass the so-called *KMP-test* [75], at least unification-based propagation is required; to eliminate intermediate data structures in general, normal-order strategy is required.

<i>transformer</i>	<i>information propagation</i>	<i>evaluation strategy</i>	<i>control restrict. variant genetic</i>		<i>KMP test</i>	<i>elim. struct.</i>
Partial evaluation	constant	in-out	poly	mono	–	–
Deforestation	constant	out-in	poly	poly	–	+
Partial deduction	unification	unspecified	poly	mono	+	–
Positive SCP	unification	out-in	poly	poly	+	+
Perfect SCP	constraint	out-in	poly	poly	+	+
GPC	constraint	out-in	poly	poly	+	+

Fig. 15. A taxonomy of transformers.

5 Larger Perspectives of Supercompilation

Supercompilation achieves program specialization, but is not limited to this application: it is a much wider framework for equivalence transformation of programs. Program inversion is one of the more advanced applications of supercompilation which we will outline in this section.

We refer to any process of simulating, analyzing or transforming programs by means of programs as *metacomputation*; the term stresses the fact that this activity is one level higher than ordinary computation (“programs as data objects”). Program specialization, composition, and inversion are different metacomputation tasks; programs that carry out these tasks, are *metaprograms*. The step from a program to the application of a metaprogram to the encoded form of the program is a *metasystem transition*; repeated use of metasystem transition leads to a *multi-level metasystem hierarchy*. In the remainder of this paper we adopt the language-independent formalization of [30] based on [27, 35, 84, 97].

Metasystem transition is a key ingredient of Turchin’s approach: the construction of hierarchies of metasystems (*e.g.* supercompilers) was taken as the basis for program analysis and transformation [83]. The book [96] defined all three Futamura projections in terms of metasystem transition.

Section 5.1 introduces a formalism for metacomputation, in Sect. 5.2 discusses supercompilation and program inversion, and Sect. 5.3 presents metasystem transition.

5.1 Metacomputation Revisited

Computation. We assume a fixed set D containing programs written in different languages, as well as their input and output data. To express the application of programs to data we define an *application language* A by the grammar

$$A ::= D \mid \langle A \ A^* \rangle$$

where the symbols $\langle, \rangle \notin D$ denote the application of a program to its inputs. Capitalized names in **typewriter** font denote arbitrary elements of D . They are free variables of the meta-notation in which the paper is written. For instance, the intended meaning of the A -expression $\langle P \ X \ Y \rangle$ is the application of program $P \in D$ to the input $X, Y \in D$. We are not interested in a specific programming language for writing programs. For simplicity, let all source-, target- and metalanguages be identical.

We write $a \Rightarrow D$ to denote the *computation* of an expression $a \in A$ to $D \in D$. For instance, $\langle P \ X \ Y \rangle \Rightarrow \text{OUT}$ is the computation of program $P \in D$ with input $X, Y \in D$ and output $\text{OUT} \in D$. Two A -expressions $a, b \in A$ are computationally equal if they can be reduced to identical D -expressions:

$$a = b \text{ iff } \forall X \in D: (a \Rightarrow X \text{ iff } b \Rightarrow X)$$

Abstraction. To represent sets of A -expressions, we define a *metacomputation language* B by the grammar

$$B ::= D \mid M \mid \langle B \ B^* \rangle$$

where M is a set of *metavariables*. A metavariable $m \in M$ is a placeholder that stands for an unspecified data element $D \in D$. We use lowercase names in **typewriter** font to denote elements of M . A B -expression b is an abstraction

that represents the set of all A -expressions obtained by replacing metavariables $\mathbf{m} \in M$ by elements of D . We write $a \in b$ to denote that $a \in A$ is an element of the set represented by $b \in B$. We refer to a B -expression also as a *configuration*.

Encoding. Expressions in the metacomputation language need to be represented as data in order to manipulate them by means of programs (ordinary computation cannot reduce B -expressions because metavariables are not in A). A *metacoding* [79] is an injective mapping $B \rightarrow D$ to encode B -expressions in D . We are not interested in a specific way of metacoding and assume some metacoding $\bar{\bullet} : B \rightarrow D$. Repeated metacoding is well-defined because $D \subset B$.

Metacomputation. It follows from our notation that $\langle \mathbf{MC} \bar{b} \rangle \Rightarrow \mathbf{D}$ denotes metacomputation on an expression $b \in B$ using a metaprogram $\mathbf{MC} \in D$. The application of \mathbf{MC} to the metacoded B -expression is an A -expression that can be reduced by ordinary computation. We should stress that this characterization of metacomputation says nothing about its concrete nature, except that it involves a metaprogram \mathbf{MC} that operates on a metacoded configuration \bar{b} . Different metaprograms may perform different operations on b , such as program specialization, program composition, or program inversion.

Definition 9 (program specializer). A program $\mathbf{SPEC} \in D$ is a *specializer* if for every program $\mathbf{P} \in D$, every input $\mathbf{X}, \mathbf{Y} \in D$ and metavariable $\mathbf{y} \in M$, there exists a program $\mathbf{R} \in D$ such that

$$\langle \mathbf{SPEC} \overline{\langle \mathbf{P} \mathbf{X} \mathbf{y} \rangle} \rangle \Rightarrow \mathbf{R} \quad \text{and} \quad \langle \mathbf{P} \mathbf{X} \mathbf{Y} \rangle = \langle \mathbf{R} \mathbf{Y} \rangle$$

Definition 10 (program composer). A program $\mathbf{CPO} \in D$ is a *composer* if for every program $\mathbf{P}, \mathbf{Q} \in D$, every input $\mathbf{X}, \mathbf{Y} \in D$ and metavariables $\mathbf{x}, \mathbf{y} \in M$, there exists a program $\mathbf{R} \in D$ such that

$$\langle \mathbf{CPO} \overline{\langle \mathbf{P} \langle \mathbf{Q} \mathbf{x} \rangle \mathbf{y} \rangle} \rangle \Rightarrow \mathbf{R} \quad \text{and} \quad \langle \mathbf{P} \langle \mathbf{Q} \mathbf{X} \rangle \mathbf{Y} \rangle = \langle \mathbf{R} \mathbf{X} \mathbf{Y} \rangle$$

Definition 11 (program inverter). A program $\mathbf{INV} \in D$ is a *program inverter* if for every program $\mathbf{P} \in D$ injective in its first argument,¹ every input $\mathbf{X}, \mathbf{Y} \in D$ and metavariable $\mathbf{x} \in M$, there exists a program $\mathbf{P}^{-1} \in D$ such that

$$\langle \mathbf{INV} \overline{\langle \mathbf{P} \mathbf{x} \mathbf{Y} \rangle} \rangle \Rightarrow \mathbf{P}^{-1} \quad \text{and} \quad \langle \mathbf{P}^{-1} \langle \mathbf{P} \mathbf{X} \mathbf{Y} \rangle \rangle \Rightarrow \mathbf{X}$$

In general when \mathbf{P} is not injective, \mathbf{P}^{-1} must return a list of answers.²

¹ \mathbf{P} is injective in its first argument if for all $\mathbf{X1}, \mathbf{X2}, \mathbf{Y} \in D$: $\langle \mathbf{P} \mathbf{X1} \mathbf{Y} \rangle = \langle \mathbf{P} \mathbf{X2} \mathbf{Y} \rangle$ implies that $\mathbf{X1}$ and $\mathbf{X2}$ are the same element of D .

² There are two types of inversion: either we are interested in an *existential* solution (one of the possible answers), or in a *universal* solution (all possible answers).

5.2 Inverse Computation by Supercompilation

A supercompiler can be used for program specialization and the transformed program in Fig. 13 illustrates a case of program composition. It is less known that supercompilation is capable of *inverse computation* [81, 65, 1] (we show later how metasystem transition can be used to generate an inverse program P^{-1}). The formulation of inverse computation is as follows. Let **EQ** be a program that tests the equality of two data elements. Given **Y**, **Z** find an **X** such that

$$\langle \text{EQ} \langle P \ X \ Y \rangle \ Z \rangle \Rightarrow \text{'True'}$$

where **'True'** is some distinct element of D . Supercompilation, more specifically driving, can be used to obtain a program **ANSWER** with answers for **x** internalized:

$$\langle \text{DRIVE} \overline{\langle \text{EQ} \langle P \ x \ Y \rangle \ Z \rangle} \rangle \Rightarrow \text{ANSWER}$$

Example 1. Let numbers be represented by lists of length n . Then program **append** a (Fig. 2) implements the addition of two numbers. Using driving (Sect. 2) we can compute $zs - ys$ by inverse computation of addition. The result of interpretive inversion for $zs = 1$ and $ys = 0$, *i.e.* driving $eq \ (a \ xs \ []) \ [1]$, appears in Fig. 16. The answer, $xs = 1$, can be extracted mechanically from the program.

$g_1 \ []$	$\triangleq \text{False}$	$g_2 \ []$	$\triangleq \text{True}$
$g_1 \ (x:xs)$	$\triangleq g_2 \ xs$	$g_2 \ (x:xs)$	$\triangleq \text{False}$

Fig. 16. Result of driving $eq \ (a \ xs \ []) \ [1]$.

Example 2. Using supercompilation instead of driving one may produce a finite program even when the list of possible answers is infinite. This may be used for theorem proving [86, 99]. An example is shown in Fig. 17 where the supercompiler (Sect. 3) is applied to $eq \ (a \ xs \ []) \ xs$ which represents the proposition $\forall n. (n + 0 = n)$ which can be proven only by using induction. The residual program constructed returns *True* for all lists. This proves the theorem.

$g_1 \ []$	$\triangleq \text{True}$
$g_1 \ (x:xs)$	$\triangleq g_1 \ xs$

Fig. 17. Result of supercompiling $eq \ (a \ xs \ []) \ xs$.

An early result for inverse computation by driving were obtained in 1972 by performing subtraction by inverse computation of binary addition [80]. In 1973 S.A. Romanenko and later S.M. Abramov implemented an algorithm, Universal Resolving Algorithm (**URA**), in which driving was combined with a mechanical extraction of answers [1, 65]. For program inversion see also [39, 41, 66, 93, 60].

In logic programming, one defines a predicate by a program $\langle P \ x \ y \rangle$ and solves the inversion problem for $Z = \text{'True'}$. Theorem proving and program

transformation are indistinguishable in the approach outlined above; they are two applications of the same equivalence transformation. The definition of predicates may be perceived as non-procedural, but their semantics is still defined in terms of computation. The application of supercompilation to problem solving and theorem proving has been discussed in [86, 87], the connection to logic programming in [1, 28, 38].

5.3 Metasystem Transition

Having introduced the basic concepts of metacomputation, we now consider the use of multi-level metasystem hierarchies together with a supercompiler. During the construction of multi-level hierarchies, we will frequently need to replace metacoded subexpressions by metavariables. The correct treatment of metacode is so essential in self-application [27], that we make elevated metavariables [97] an integral part of the MST-language. We define a *metasystem transition language* C by the grammar

$$C ::= D \mid M_N \mid \langle C \ C^* \rangle$$

where M_N is a set of *elevated metavariables* \mathbf{m}_H , $H \in N$. An elevated metavariable \mathbf{m}_H ranges over data metacoded H -times. We will denote by D^H the set of metacode $\overline{\mathbf{d}}^H$ of all $\mathbf{d} \in D$. A metavariable without elevation has $\mathbf{0}$ as its elevation index. A C -expression c represents the set of all A -expressions obtained by substituting elevated metavariables \mathbf{m}_H by elements of D^H .

Metasystem Transition. The construction of each next level in a metasystem hierarchy, *i.e.* each *metasystem transition* (MST) [96], is done in three steps [36]:

- (A) given an initial A -expression a , ('computation')
- (B) define a C -expression c such that $a \in c$, ('abstraction')
- (C) apply a metaprogram \mathbf{MC} to the metacode \overline{c} . ('metacomputation')

The expression obtained in the last step is again an A -expression and the same procedure can be repeated. Expressions obtained by MST are called *MST-formulas*. This definition says nothing about the goal of the MST, except that it is an abstraction of an A -expression a to a configuration c , followed by the application of a metaprogram \mathbf{MC} to \overline{c} .

Generating Inverse Programs. The inverse computation of a program can always be performed using driving, but the performance can be poor whilst often more efficient inverse programs are known to exist. Figure 18 show how MST can be used to produce inverse programs by specialization of the universal resolving algorithm **URA** [1]; see [66]. For notational convenience let $\langle \mathbf{Q} \ \mathbf{x} \ \mathbf{y} \ \mathbf{z} \rangle$ be defined by $\langle \mathbf{EQ} \ \langle \mathbf{P} \ \mathbf{x} \ \mathbf{y} \rangle \ \mathbf{z} \rangle$. A specializer **SPEC** is used for the sake of generality, but it should be clear that a supercompiler **SCP** can be used instead.

1st MST Define a C -expression (B0) by replacing \mathbf{X} by \mathbf{x}_0 in the A -expression (A0), and apply **URA** to the metacoded C -expression (A1) to perform inverse computation: the 1st MST. Inverse computation of \mathbf{Q} is achieved.

2nd MST Define a C -expression (B1) by replacing \bar{Y}, \bar{Z} by y_1, z_1 in the A -expression (A1)³, and apply **SPEC'** to the metacoded C -expression (A2) to specialize **URA** and remove its interpretive overhead: the 2nd MST. The result is an inverted program Q^{-1} that returns **ANSWER** given Y, Z .

3rd MST Define a C -expression (B2) by replacing \bar{q} by q_2 in the A -expression (A2), and apply SPEC'' to the metacoded C -expression (A3): the 3rd MST. The result is an inverter **INV** that converts a program q into q^{-1} .

4th MST Define a *C*-expression (B3) by replacing $\overline{\overline{\text{URA}}}$ by ura_2 in the *A*-expression (A3), and apply SPEC' to the metacoded *C*-expression (A4): the 4th MST. The result is an inverter generator **INVGEN**.

(A0)	$\langle Q \ X \ Y \ Z \rangle \Rightarrow \text{B00L}$	(computation)
(B0)	$\langle Q \ x_0 \ Y \ Z \rangle$	(abstraction)
.....		
(A1)	$\langle \text{URA} \ \overline{\langle Q \ x_0 \ Y \ Z \rangle} \rangle \Rightarrow \text{ANSWER}$	(1st MST)
(B1)	$\langle \text{URA} \ \overline{\langle Q \ x_0 \ y_1 \ z_1 \rangle} \rangle$	(abstraction)
.....		
(A2)	$\langle \text{SPEC}' \ \overline{\langle \text{URA} \ \overline{\langle Q \ x_0 \ y_1 \ z_1 \rangle} \rangle} \rangle \Rightarrow Q^{-1}$	(2nd MST)
(B2)	$\langle \text{SPEC}' \ \overline{\langle \text{URA} \ \overline{\langle q_2 \ x_0 \ y_1 \ z_1 \rangle} \rangle} \rangle$	(abstraction)
.....		
(A3)	$\langle \text{SPEC}'' \ \overline{\langle \text{SPEC}' \ \overline{\langle \text{URA} \ \overline{\langle q_2 \ x_0 \ y_1 \ z_1 \rangle} \rangle} \rangle} \rangle \Rightarrow \text{INV}$	(3rd MST)
(B3)	$\langle \text{SPEC}'' \ \overline{\langle \text{SPEC}' \ \overline{\langle \text{ura}_2 \ \overline{\langle q_2 \ x_0 \ y_1 \ z_1 \rangle} \rangle} \rangle} \rangle$	(abstraction)
.....		
(A4)	$\langle \text{SPEC}''' \ \overline{\langle \text{SPEC}'' \ \overline{\langle \text{SPEC}' \ \overline{\langle \text{ura}_2 \ \overline{\langle q_2 \ x_0 \ y_1 \ z_1 \rangle} \rangle} \rangle} \rangle} \rangle \Rightarrow \text{INVGEN}$	(4th MST)

Fig. 18. MST-formulas for program inversion.

A hierarchy of metasystems can be visualized using a *2-dimensional notation*⁴: (i) an expression is moved down one line down for each metacoding; (ii) the elevation of a metavariable \mathbf{m}_H is shown by a bullet \bullet located H lines below the metavariable. Consider the two-dimensional version of the 3rd MST (A3) as example:

$$\begin{array}{c} \langle \text{SPEC}'' \rangle \xrightarrow{\hspace{10em}} \Rightarrow \text{INV} \\ \langle \text{SPEC}' \xrightarrow{\text{q}} \hspace{1.5em} \rangle \\ \langle \text{URA} \mid \text{y} \text{ z} \rangle \\ \langle \bullet \text{ x } \bullet \bullet \rangle \end{array}$$

The *Futamura projections* [23] are, in all probability, the first example of program transformation beyond a single metasystem level, but MST is not lim-

³ We take the liberty to replace subexpressions of $d \in D$ by metavariables and interrupt the horizontal line above the enclosing expression; defined formally in [30].

⁴ Introduced by Turchin; a preliminary form appeared in [27].

ited to this application. Turchin suggested MST to increase the power of theorem proving [86] and it inspired a constructive approach to the foundations of mathematics [88]. The philosophical background of MST was exposed in [82]; see also [91, 92].

The first successful self-application of a partial evaluator was reported in [45] and several self-application partial evaluators have been built since then. The book [44] describes approaches to and systems for partial evaluation, and includes a comprehensive bibliography on this topic. A self-applicable partial evaluator for a functional language with Refal data structures is described in [68]. The generation of an algorithm representing binary subtraction from binary addition by self-application of a simple supercompiler was reported in [39]. A self-applicable supercompiler for Refal is described in [98, 60]. Examples of MST include multiple self-application [27, 32], the generation of program transformers [29, 31], the implementation of non-standard semantics [3], and other related work on metacomputation and MST [1, 2, 28, 36, 37, 42, 49, 48].

Acknowledgments. Thanks to Sergei Abramov, Andrei Klimov, Andrei Nemytykh, Victoria Pinchuk, Alexander Romanenko (†), Sergei Romanenko, and last but not least, Valentin F. Turchin for many stimulating discussions and hospitality during visits in Russia and New York. Special thanks are due to Andrei Klimov for transcription and translation of Russian titles. Thanks to John Hatcliff, Neil D. Jones, Jesper Jørgensen, Bern Martens, Kristian Nielsen, and David Sands for stimulating discussions on various topics of this paper. Finally, we are indebted to all members of the Topps group at DIKU for providing an excellent working environment.

The first author was partially supported by an Erwin-Schrödinger-Fellowship of the Austrian Science Foundation (FWF) under grant J0780 & J0964. Both authors were also supported by the DART project funded by the Danish Natural Sciences Research Council.

References

1. S.M. Abramov. Metavychislenija i logicheskoe programmirovanie (Metacomputation and logic programming). *Programmirovanie*, 3:31–44, 1991. (In Russian).
2. S.M. Abramov. Metacomputation and program testing. In *1st International Workshop on Automated and Algorithmic Debugging*, pp. 121–135, 1993.
3. S.M. Abramov. *Metavychislenija i ikh prilozhenija (Metacomputation and its application)*. Nauka, Moscow, 1995. (In Russian).
4. S.M. Abramov and N.V. Kondratiev. Kompiljator, osnovannyj na metode chas-tichnykh vychislenij (A compiler based on the method of partial evaluation). In *Nekotorye voprosy prikladnoj matematiki i programmogo obespechenija EhVB*, pp. 66–69. Moscow State University, Moscow, 1982. (in Russian).
5. M. Alpuente, M. Falaschi, and G. Vidal. Narrowing-driven partial evaluation of functional logic programs. In *European Symposium on Programming—ESOP ’96*, LNCS. Springer-Verlag, 1996. To appear.
6. L. Beckman, A. Haraldson, Ö. Oskarsson, and E. Sandewall. A partial evaluator and its use as a programming tool. *Artificial Intelligence*, 7:319–357, 1976.
7. R. Bird and P.L. Wadler. *Introduction to Functional Programming*. Prentice-Hall, 1988.

8. D. Bjørner, A.P. Ershov, and N.D. Jones, editors. *Partial Evaluation and Mixed Computation*. North-Holland, Amsterdam, 1988.
9. R. Bol. Loop checking in partial deduction. *Journal of Logic Programming*, 16(1&2):25–46, 1993.
10. M.A. Bulyonkov. Polyvariant mixed computation for analyzer programs. *Acta Informatica*, 21:473–484, 1984.
11. R.M. Burstall and J. Darlington. A transformation system for developing recursive programs. *Journal of the Association for Computing Machines*, 24(1):44–67, 1977.
12. C.-L. Chang and R.C.-T. Lee. *Symbolic Logic and Mechanical Theorem Proving*. Computer Science and Applied Mathematics. Academic Press, 1973.
13. C. Consel and O. Danvy. For a better support of static data flow. In J. Hughes, editor, *Functional Programming and Computer Architecture*, vol. 523 of *LNCS*, pp. 495–519. Springer-Verlag, 1991.
14. C. Consel and O. Danvy. Tutorial notes on partial evaluation. In *ACM Symposium in Principles of Programming Languages*, pages 493–501. ACM Press, 1993.
15. C. Consel and S.C. Khoo. Parameterized partial evaluation. *ACM TOPLAS*, 15(3):463–493, 1993.
16. O. Danvy, R. Glück, and P. Thiemann, editors. *Partial Evaluation. Proceedings*. LNCS. Springer-Verlag, 1996. To appear.
17. N. Dershowitz and J.-P. Jouannaud. Rewrite systems. In J. van Leeuwen, editor, *Handbook of Theoretical Computer Science*, pp. 244–320. Elsevier, 1992.
18. J. Dixon. The specializer, a method of automatically writing computer programs. Technical report, Division of Computer Research and Technology, National Institute of Health, Bethesda, Maryland, 1971.
19. A.P. Ershov. On the partial computation principle. *Information Processing Letters*, 6(2):38–41, 1977.
20. A.P. Ershov. On the essence of compilation. In E.J. Neuhold, editor, *Formal Description of Programming Concepts*, pp. 391–420. North-Holland, 1978.
21. A. Ferguson and P.L. Wadler. When will deforestation stop? In *1988 Glasgow Workshop on Functional Programming*, pages 39–56, 1988.
22. S.N. Florencev, Y.V. Oljunin, and V.F. Turchin. (An efficient interpreter for the language Refal). Preprint, Institute of Applied Mathematics, Academy of Sciences of the USSR, Moscow, 1969. (in Russian).
23. Y. Futamura. Partial evaluation of computing process – an approach to a compiler-compiler. *Systems, Computers, Controls*, 2(5):45–50, 1971.
24. Y. Futamura. Partial computation of programs. In E. Goto, K. Furukawa, R. Nakajima, I. Nakata, and A. Yonezawa, editors, *RIMS Symposia on Software Science and Engineering*, vol. 147 of *LNCS*, pp. 1–35, Kyoto, Japan, 1983. Springer-Verlag.
25. Y. Futamura. Program evaluation and generalized partial computation. In *International Conference on Fifth Generation Computer Systems*, pp. 1–8, 1988.
26. J. Gallagher. Tutorial in specialisation of logic programs. In *Symposium on Partial Evaluation and Semantics-Based Program Manipulation*, pp. 88–98. ACM Press, 1993.
27. R. Glück. Towards multiple self-application. In *Proceedings of the Symposium on Partial Evaluation and Semantics-Based Program Manipulation*, pp. 309–320. ACM Press, 1991.
28. R. Glück. Projections for knowledge based systems. In R. Trappl, editor, *Cybernetics and Systems Research'92*, volume 1, pp. 535–542. World Scientific, 1992.

29. R. Glück. On the generation of specializers. *Journal of Functional Programming*, 4(4):499–514, 1994.
30. R. Glück. On the mechanics of metasystem hierarchies in program transformation. In M. Proietti, editor, *Logic Program Synthesis and Transformation. Proceedings*, vol. 1048 of *LNCS*, pp. 234–251. Springer-Verlag, 1996.
31. R. Glück and J. Jørgensen. Generating transformers for deforestation and supercompilation. In B. Le Charlier, editor, *Static Analysis. Proceedings*, volume 864 of *LNCS*, pp. 432–448, Namur, Belgium, 1994. Springer-Verlag.
32. R. Glück and J. Jørgensen. Efficient multi-level generating extensions for program specialization. In S.D. Swierstra and M. Hermenegildo, editors, *Programming Languages, Implementations, Logics and Programs (PLILP'95)*, vol. 982 of *LNCS*, pp. 259–278. Springer-Verlag, 1995.
33. R. Glück, J. Jørgensen, B. Martens, and M.H. Sørensen. Controlling conjunctive partial deduction of definite logic programs. Technical Report CW 226, Katholieke Universiteit Leuven, 1996.
34. R. Glück and A.V. Klimov. Occam's razor in metacomputation: the notion of a perfect process tree. In P. Cousot, M. Falaschi, G. Filè, and G. Rauzy, editors, *Static Analysis. Proceedings.*, vol. 724 of *LNCS*, pp. 112–123. Springer-Verlag, 1993.
35. R. Glück and A.V. Klimov. Metacomputation as a tool for formal linguistic modeling. In R. Trappl, editor, *Cybernetics and Systems'94*, volume 2, pp. 1563–1570. World Scientific, 1994.
36. R. Glück and A.V. Klimov. Metasystem transition schemes in computer science and mathematics. *World Futures*, 45:213–243, 1995.
37. R. Glück and A.V. Klimov. Reduction of language hierarchies. In *Proceedings of the 14th International Congress on Cybernetics*, page To appear, Namur, Belgium, 1995. International Association for Cybernetics.
38. R. Glück and M.H. Sørensen. Partial deduction and driving are equivalent. In M. Hermenegildo and J. Penjam, editors, *Programming Language Implementation and Logic Programming. Proceedings*, vol. 844 of *LNCS*, pp. 165–181. Springer-Verlag, 1994.
39. R. Glück and V.F. Turchin. Application of metasystem transition to function inversion and transformation. In *Proceedings of the ISSAC'90 (Tokyo, Japan)*, pp. 286–287. ACM Press, 1990.
40. R. Gurin and S.A. Romanenko. *Jazyk programirovaniya Refal Plus (The Refal Plus programming language)*. Intertech, Moscow, 1991. (in Russian).
41. P.G. Harrison. Function inversion. In Bjørner et al. [8], pp. 153–166.
42. J. Hatcliff and Robert Glück. Reasoning about hierarchies of online program specialization systems. In Danvy et al. [16]. To appear.
43. N.D. Jones. The essence of program transformation by partial evaluation and driving. In N.D. Jones, M. Hagiya, and M. Sato, editors, *Logic, Language, and Computation*, vol. 792 of *LNCS*, pp. 206–224. Springer-Verlag, 1994. Festschrift in honor of S.Takasu.
44. N.D. Jones, C.K. Gomard, and P. Sestoft. *Partial Evaluation and Automatic Program Generation*. Prentice-Hall, 1993.
45. N.D. Jones, P. Sestoft, and H. Søndergaard. An experiment in partial evaluation: the generation of a compiler generator. In J.-P. Jouannaud, editor, *Rewriting Techniques and Applications, Dijon, France.*, vol. 202 of *LNCS*, pp. 124–140. Springer-Verlag, 1985.

46. N.D. Jones, P. Sestoft, and H. Søndergaard. Mix: a self-applicable partial evaluator for experiments in compiler generation. *Lisp and Symbolic Computation*, 2(1):9–50, 1989.
47. S.C. Kleene. *Introduction to Metamathematics*. Van Nostrand, 1952.
48. Andrei V. Klimov. Dynamic specialization in extended functional language with monotone objects. In *Proceedings of the Symposium on Partial Evaluation and Semantics-Based Program Manipulation*, pp. 199–210. ACM Press, 1991.
49. A.V. Klimov and S.A. Romanenko. Metavychislitel' dlja jazyka Refal. Osnovnye ponjatija i primery. (A metaevaluator for the language Refal. Basic concepts and examples). Preprint 71, Keldysh Institute of Applied Mathematics, Academy of Sciences of the USSR, Moscow, 1987. (in Russian).
50. A.V. Klimov and S.A. Romanenko. Sistema programirovanija Refal-2 dlja ES. Opisanie vkhodnogo jazyka (Programming system Refal-2 for ES computers. The source language description). Technical report, Keldysh Institute of Applied Mathematics, Academy of Sciences of the USSR, Moscow, 1987. (in Russian).
51. A.V. Klimov, S.A. Romanenko, and V.F. Turchin. Teoreticheskie osnovy sintaksicheskogo otozhdestvlenija v jazyke Refal (The theory of pattern matching in Refal). Preprint 13, Keldysh Institute of Applied Mathematics, Academy of Sciences of the USSR, Moscow, 1973. (in Russian).
52. J. Komorowski. An introduction to partial deduction. In A. Pettorossi, editor, *Meta-Programming in Logic*, vol. 649 of *LNCS*, pp. 49–69, 1992.
53. J. Komorowski. Special issue on partial deduction. *Journal of Logic Programming*, 16(1&2):1–189, 1993.
54. M. Leuschel, D. De Schreye, and A. de Waal. A conceptual embedding of folding into partial deduction: Towards a maximal integration. Technical Report CW 225, Katholieke Universiteit Leuven, 1995.
55. M. Leuschel and B. Martens. Global control for partial deduction through characteristic atoms and global trees. In Danvy et al. [16]. To appear.
56. J.W. Lloyd and J.C. Shepherdson. Partial evaluation in logic programming. *Journal of Logic Programming*, 11(3-4):217–242, 1991.
57. L.A. Lombardi. Incremental computation. In F. L. Alt and M. Rubinoff, editors, *Advances in Computers*, volume 8, pp. 247–333. Academic Press, 1967.
58. L.A. Lombardi and B. Raphael. Lisp as the language for an incremental computer. In E.C. Berkeley and D.G. Bobrow, editors, *The Programming Language Lisp: Its Operation and Applications*, pp. 204–219, Cambridge, Massachusetts, 1964. MIT Press.
59. B. Martens and J. Gallagher. Ensuring global termination of partial deduction while allowing flexible polyvariance. In L. Stirling, editor, *International Conference on Logic Programming*, pp. 597–613. MIT Press, 1995.
60. A.P. Nemytykh, V.A. Pinchuk, and V.F. Turchin. A self-applicable supercompiler. In Danvy et al. [16]. To appear.
61. K. Nielsen and M.H. Sørensen. Call-by-name CPS-translation as a binding-time improvement. In A. Mycroft, editor, *Static Analysis*, vol. 983 of *LNCS*, pp. 296–313. Springer-Verlag, 1995.
62. R.M. Nirenberg. A practical turing machine representation. *SIGACT News*, 17(3):35–44, 1986.
63. M. Proietti and A. Pettorossi. Unfolding – definition – folding, in this order for avoiding unnecessary variables in logic programs. In *Programming Language Implementation and Logic Programming*, vol. 528 of *LNCS*, pp. 347–358. Springer-Verlag, 1991.

64. G. Puebla and M. Hermenegildo. Implementation of multiple specialization in logic programs. In *Proceedings of the Symposium on Partial Evaluation and Semantics-Based Program Manipulation*, pp. 77–87. ACM Press, 1995.
65. A.Y. Romanenko. The generation of inverse functions in Refal. In Bjørner et al. [8], pp. 427–444.
66. A.Y. Romanenko. Inversion and metacomputation. In *Proceedings of the Symposium on Partial Evaluation and Semantics-Based Program Manipulation. (Yale University, Connecticut)*, pages 12–22. ACM Press, 1991.
67. S.A. Romanenko. Progonka dlja programm na Refale-4 (Driving for Refal-4 programs). Preprint 211, Keldysh Institute of Applied Mathematics, Academy of Sciences of the USSR, Moscow, 1987. (in Russian).
68. S.A. Romanenko. A compiler generator produced by a self-applicable specializer can have a surprisingly natural and understandable structure. In Bjørner et al. [8], pp. 445–463.
69. S.A. Romanenko. Arity raiser and its use in program specialization. In N.D. Jones, editor, *ESOP'90*, vol. 432 of *LNCS*, pages 341–360. Springer-Verlag, 1990.
70. D. Sands. Proving the correctness of recursion-based automatic program transformation. In P. Mosses, M. Nielsen, and M.I. Schwartzbach, editors, *Theory and Practice of Software Development*, vol. 915 of *LNCS*, pages 681–695. Springer-Verlag, 1995.
71. D. Sands. Total correctness by local improvement in program transformation. In *22nd Symposium on Principles of Programming Languages*, pages 221–232. ACM Press, 1995.
72. D. Smith. Partial evaluation of pattern matching in constraint logic programming. In *Symposium on Partial Evaluation and Semantics-Based Program Manipulation*, pp. 62–71. ACM Press, 1991.
73. M.H. Sørensen. Turchin's supercompiler revisited. Master's thesis, Department of Computer Science, University of Copenhagen, 1994. DIKU-rapport 94/17.
74. M.H. Sørensen and R. Glück. An algorithm of generalization in positive supercompilation. In J.W. Lloyd, editor, *Logic Programming: Proceedings of the 1995 International Symposium*, pp. 465–479. MIT Press, 1995.
75. M.H. Sørensen, R. Glück, and N.D. Jones. Towards unifying deforestation, supercompilation, partial evaluation, and generalized partial computation. In D. Sannella, editor, *Programming Languages and Systems*, vol. 788 of *LNCS*, pp. 485–500. Springer-Verlag, 1994.
76. M.H. Sørensen, R. Glück, and N.D. Jones. A positive supercompiler. *Journal of Functional Programming*, 1996. To appear.
77. A. Takano. Generalized partial computation using disunification to solve constraints. In M. Rusinowitch and J.L. Remy, editors, *Conditional Term Rewriting Systems. Proceedings*, vol. 656 of *LNCS*, pp. 424–428. Springer-Verlag, 1993.
78. V.F. Turchin. Metajazyk dlja formal'nogo opisanija algoritmicheskikh jazykov (A metalanguage for the formal description of algorithmic languages). In *Cifrovaja Vychislitel'naja Tekhnika i Programirovanie*, pp. 116–124. Sovetskoe Radio, Moscow, 1966. (in Russian).
79. V.F. Turchin. Programirovanie na jazyke Refal. (Programming in the language Refal). Preprint 41, 43, 44, 48, 49, Institute of Applied Mathematics, Academy of Sciences of the USSR, Moscow, 1971. (in Russian).
80. V.F. Turchin. Ehkvivalentnye preobrazovaniya rekursivnykh funkciy na Refale (Equivalent transformations of recursive functions defined in Refal). In *Teorija*

- Jazykov i Metody Programirovaniya (Proceedings of the Symposium on the Theory of Languages and Programming Methods)*, pages 31–42, Kiev-Alushta, USSR, 1972. (In Russian).
81. V.F. Turchin. Ehkvivalentnye preobrazovanija programm na Refale (Equivalent transformations of Refal programs). *Avtomatizirovannaja Sistema upravlenija stroitel'stvom. Trudy CNIPIASS*, 6:36–68, 1974. (In Russian).
 82. V.F. Turchin. *The Phenomenon of Science*. Columbia University Press, New York, 1977.
 83. V.F. Turchin. A supercompiler system based on the language Refal. *SIGPLAN Notices*, 14(2):46–54, 1979.
 84. V.F. Turchin. The language Refal, the theory of compilation and metasystem analysis. Courant Computer Science Report 20, Courant Institute of Mathematical Sciences, New York University, 1980.
 85. V.F. Turchin. Semantic definitions in Refal and the automatic production of compilers. In N.D. Jones, editor, *Workshop on Semantics-Directed Compiler Generation, Århus, Denmark*, volume 94 of *LNCS*, pp. 441–474. Springer-Verlag, 1980.
 86. V.F. Turchin. The use of metasystem transition in theorem proving and program optimization. In J.W. de Bakker and J. van Leeuwen, editors, *Automata, Languages and Programming*, volume 85 of *LNCS*, pp. 645–657, Noordwijkerhout, Netherlands, 1980. Springer-Verlag.
 87. V.F. Turchin. The concept of a supercompiler. *Transactions on Programming Languages and Systems*, 8(3):292–325, 1986.
 88. V.F. Turchin. A constructive interpretation of the full set theory. *The Journal of Symbolic Logic*, 52(1):172–201, 1987.
 89. V.F. Turchin. The algorithm of generalization. In Bjørner et al. [8], pp. 531–549.
 90. V.F. Turchin. *Refal-5, Programming Guide and Reference Manual*. New England Publishing Co., Holyoke, Massachusetts, 1989.
 91. V.F. Turchin. The cybernetic ontology of action. *Kybernetes*, 22(2):10–30, 1993.
 92. V.F. Turchin. On cybernetic epistemology. *Systems Research*, 10(1):3–28, 1993.
 93. V.F. Turchin. Program transformation with metasystem transitions. *Journal of Functional Programming*, 3(3):283–313, 1993.
 94. V.F. Turchin. On generalization of lists and strings in supercompilation. Technical report, City College of the City University of New York, 1995.
 95. V.F. Turchin. Metacomputation: MST plus SCP. In Danvy et al. [16]. To appear.
 96. V.F. Turchin, And.V. Klimov, Ark.V. Klimov, V.F. Khoroshevsky, A.G. Krasovsky, S.A. Romanenko, I.B. Shchenkov, and E.V. Travkina. *Bazisnyj Refal i ego realizacija na vychislitelnykh mashinakh (Basic Refal and its implementation on computers)*. GOSSTROJ SSSR, CNIPIASS, Moscow, 1977. (in Russian).
 97. V.F. Turchin and A.P. Nemytykh. Metavariables: their implementation and use in program transformation. Technical Report CSc. TR 95-012, City College of the City University of New York, 1995.
 98. V.F. Turchin and A.P. Nemytykh. A self-applicable supercompiler. Technical Report CSc. TR 95-010, City College of the City University of New York, 1995.
 99. V.F. Turchin, R. Nirenberg, and D. Turchin. Experiments with a supercompiler. In *Conference Record of the ACM Symposium on Lisp and Functional Programming*, pp. 47–55. ACM Press, 1982.
 100. P.L. Wadler. Deforestation: Transforming programs to eliminate intermediate trees. *Theoretical Computer Science*, 73:231–248, 1990. Preliminary version in ESOP'88 LNCS vol. 300.

This article was processed using the L^AT_EX macro package with LLNCS style